

WPI

Lecture 12 - Vision

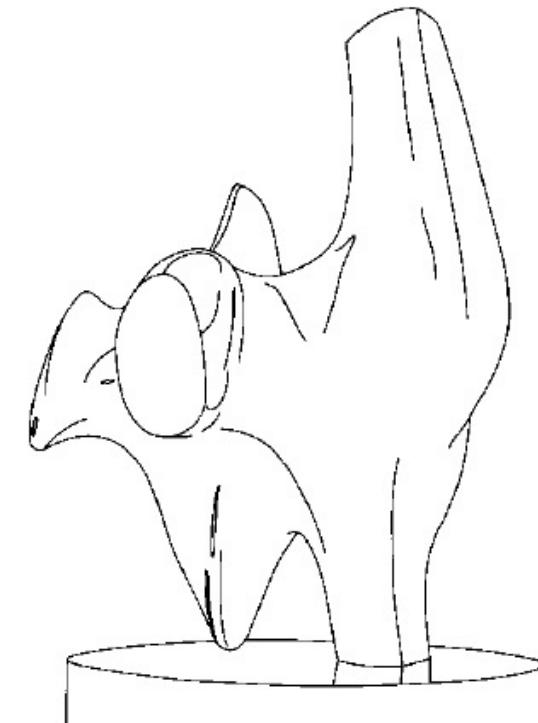
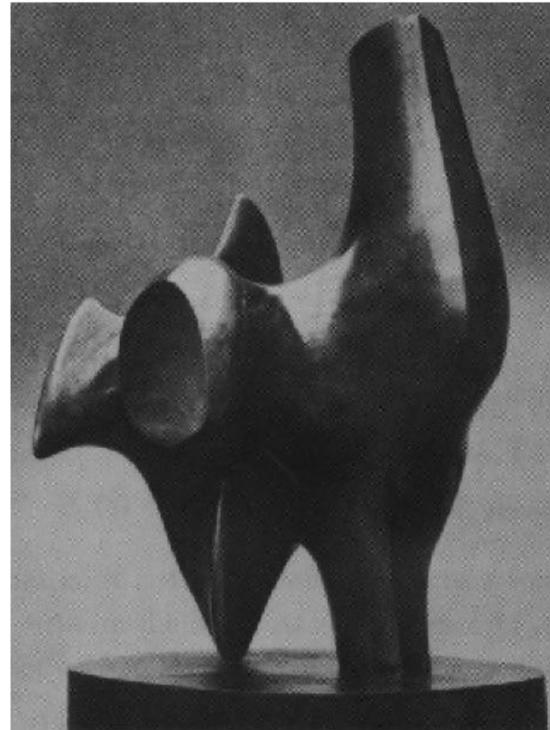
Image gradients and edges

Feature Extraction and Detection

Alexandros Lioulemes, PhD

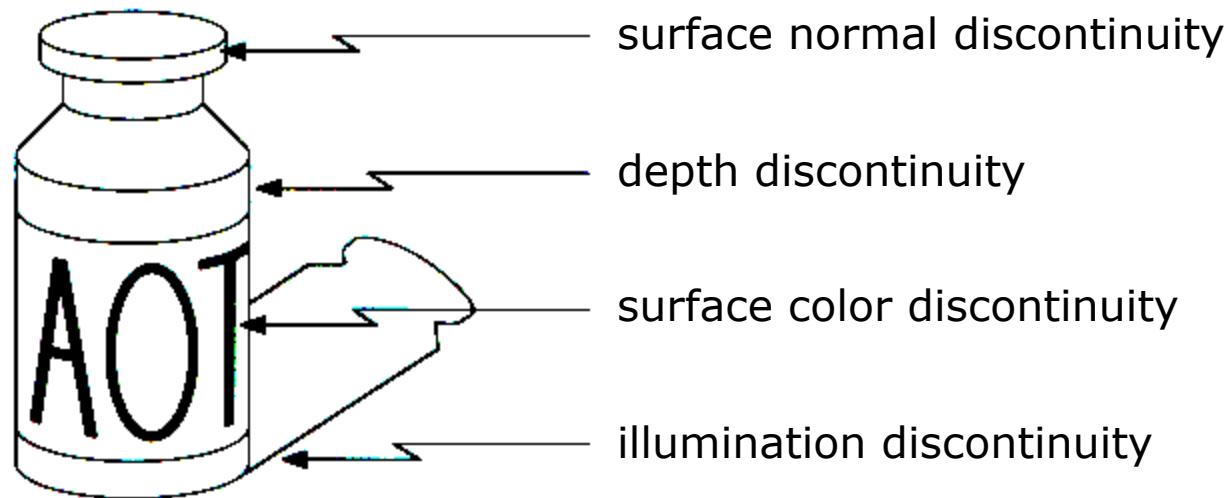


Edge detection



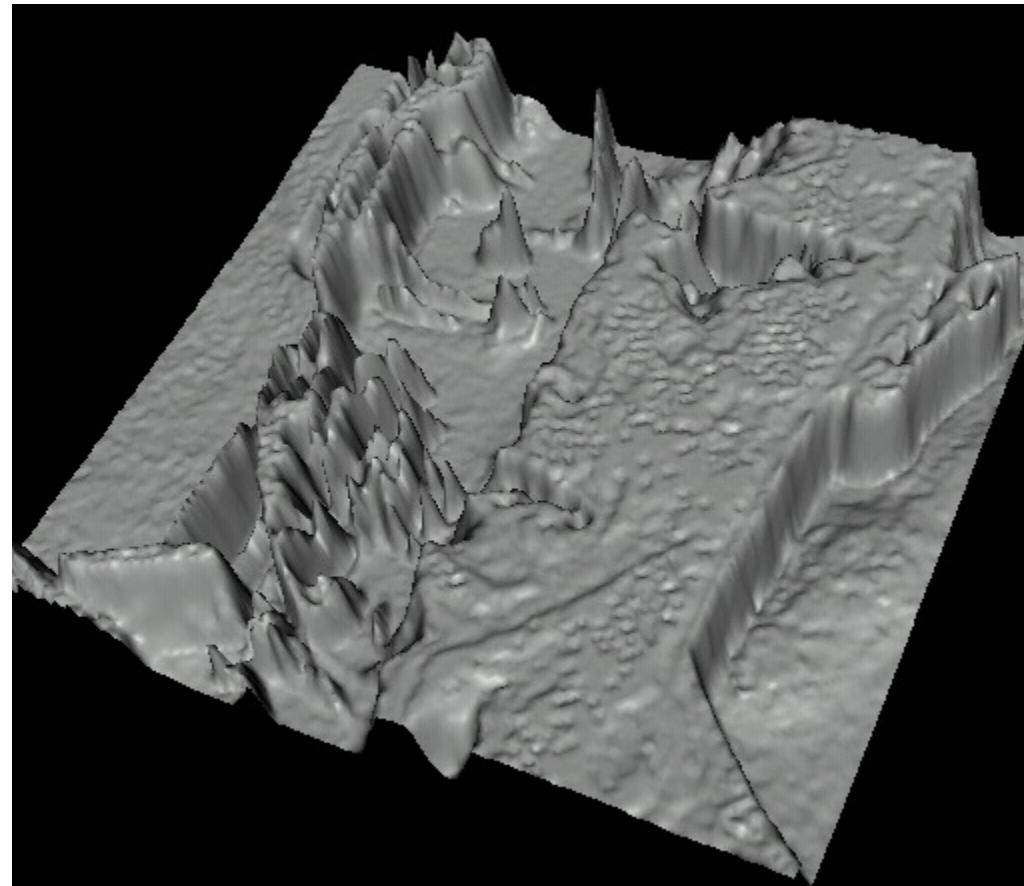
- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels

Origin of edges



- Edges are caused by a variety of factors

Images as function



- Edges look like steep cliffs

Characterizing edges

- An edge is a place of *rapid change* in the image intensity function

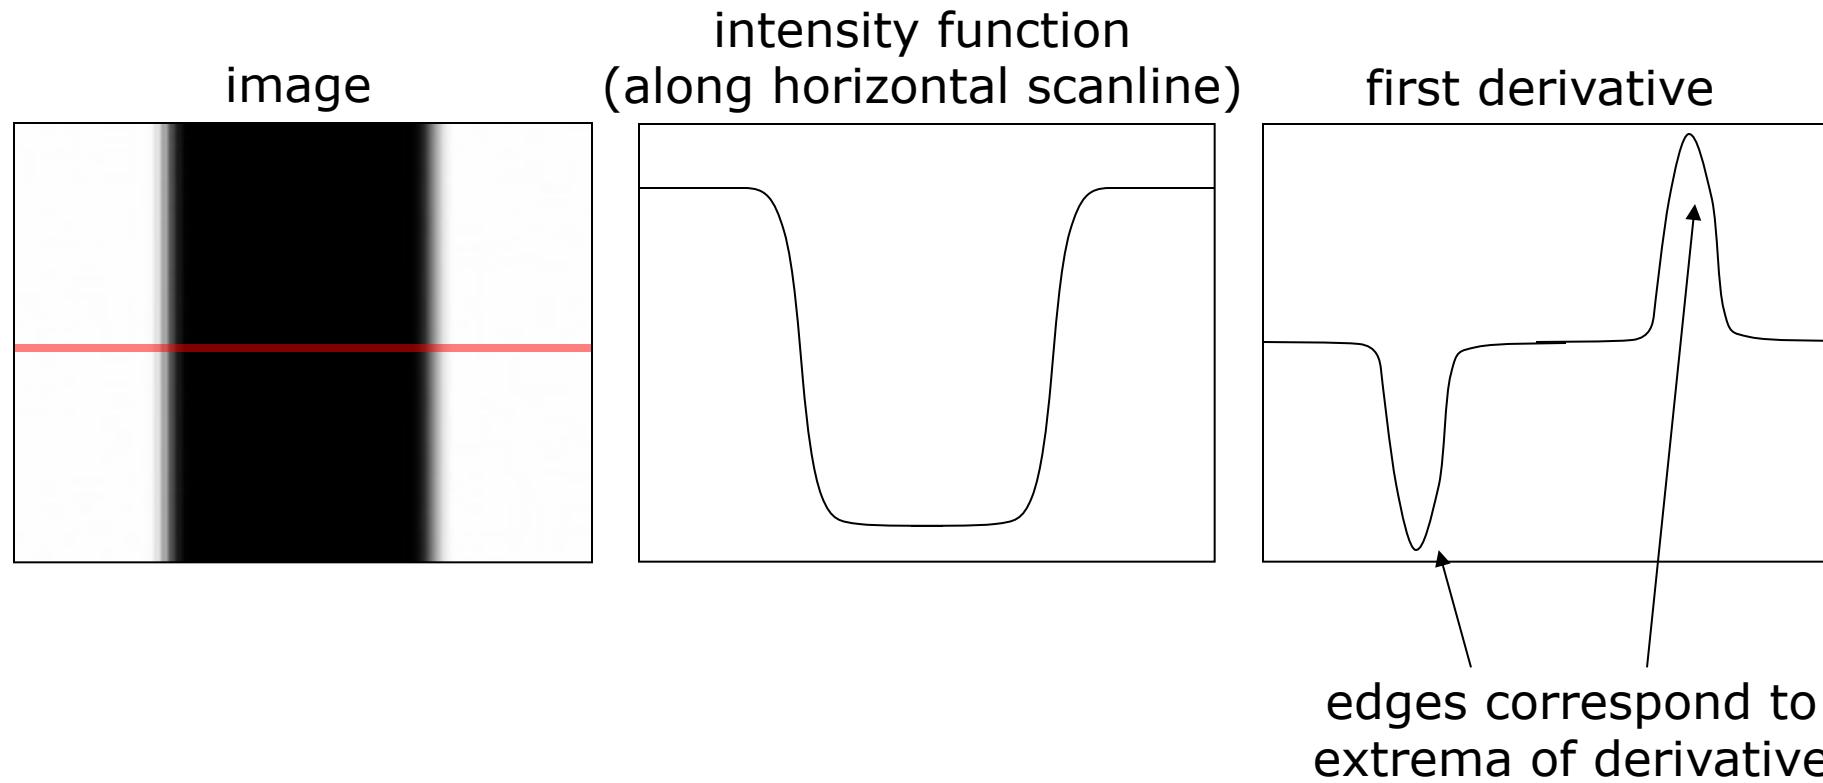


Image derivatives

- How can we differentiate a *digital* image $F[x,y]$?
 - Option 1: reconstruct a continuous image, f , then compute the derivative
 - Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x} : \begin{array}{|c|c|c|}\hline & & \\ \hline 1 & -1 & \\ \hline & & \\ \hline \end{array}$$

H_x

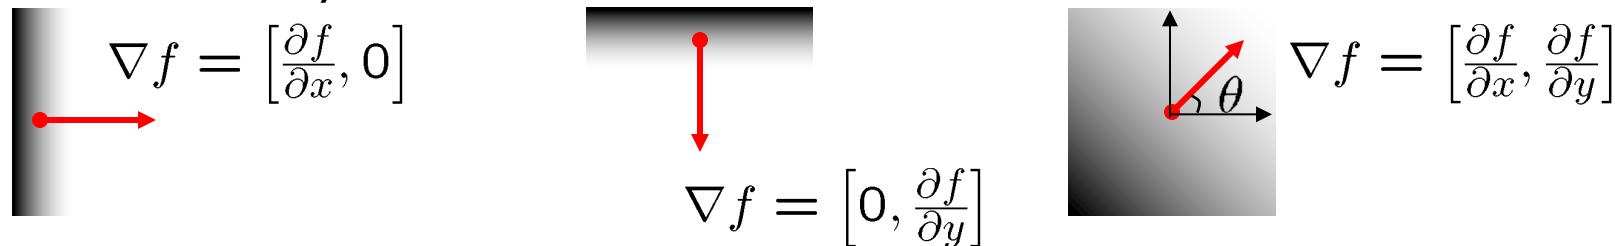
$$\frac{\partial f}{\partial y} : \begin{array}{|c|c|c|}\hline & & \\ \hline & & -1 \\ \hline & & 1 \\ \hline \end{array}$$

H_y

Image gradient

- The *gradient* of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The *edge strength* is given by the gradient magnitude:

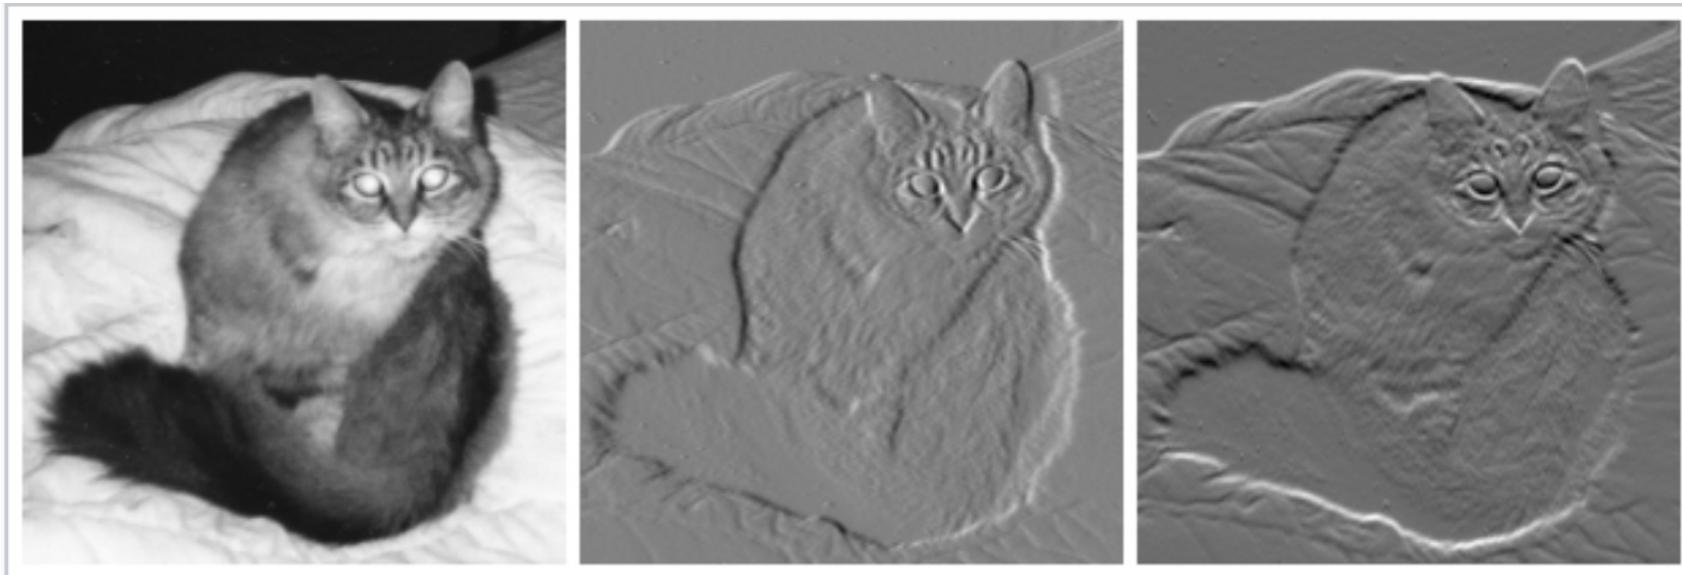
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:

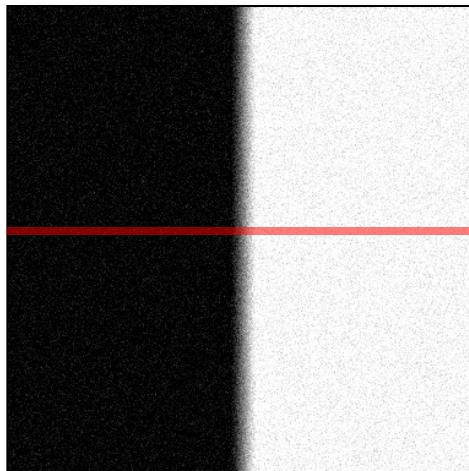
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?

Image gradient

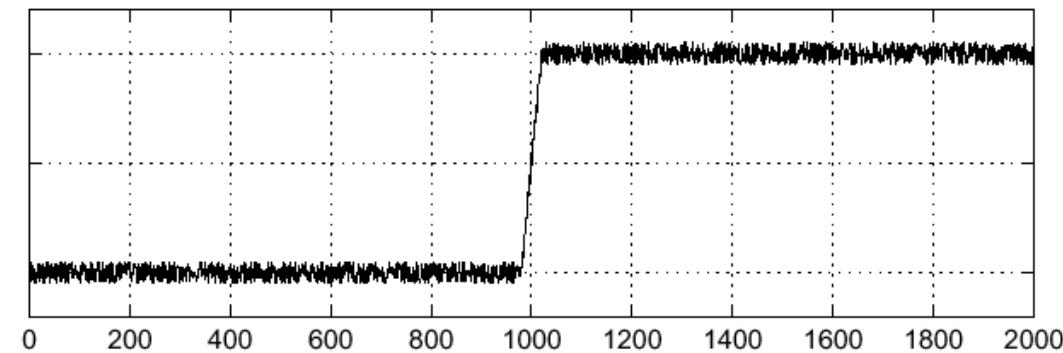


Effects of noise

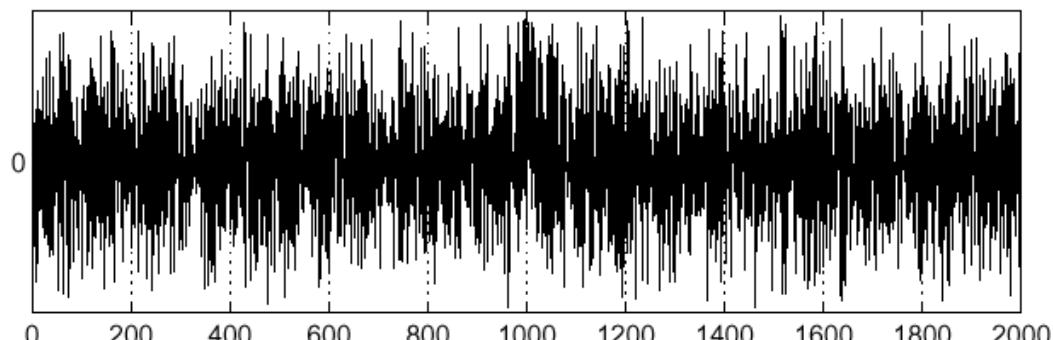


Noisy input image

$$f(x)$$

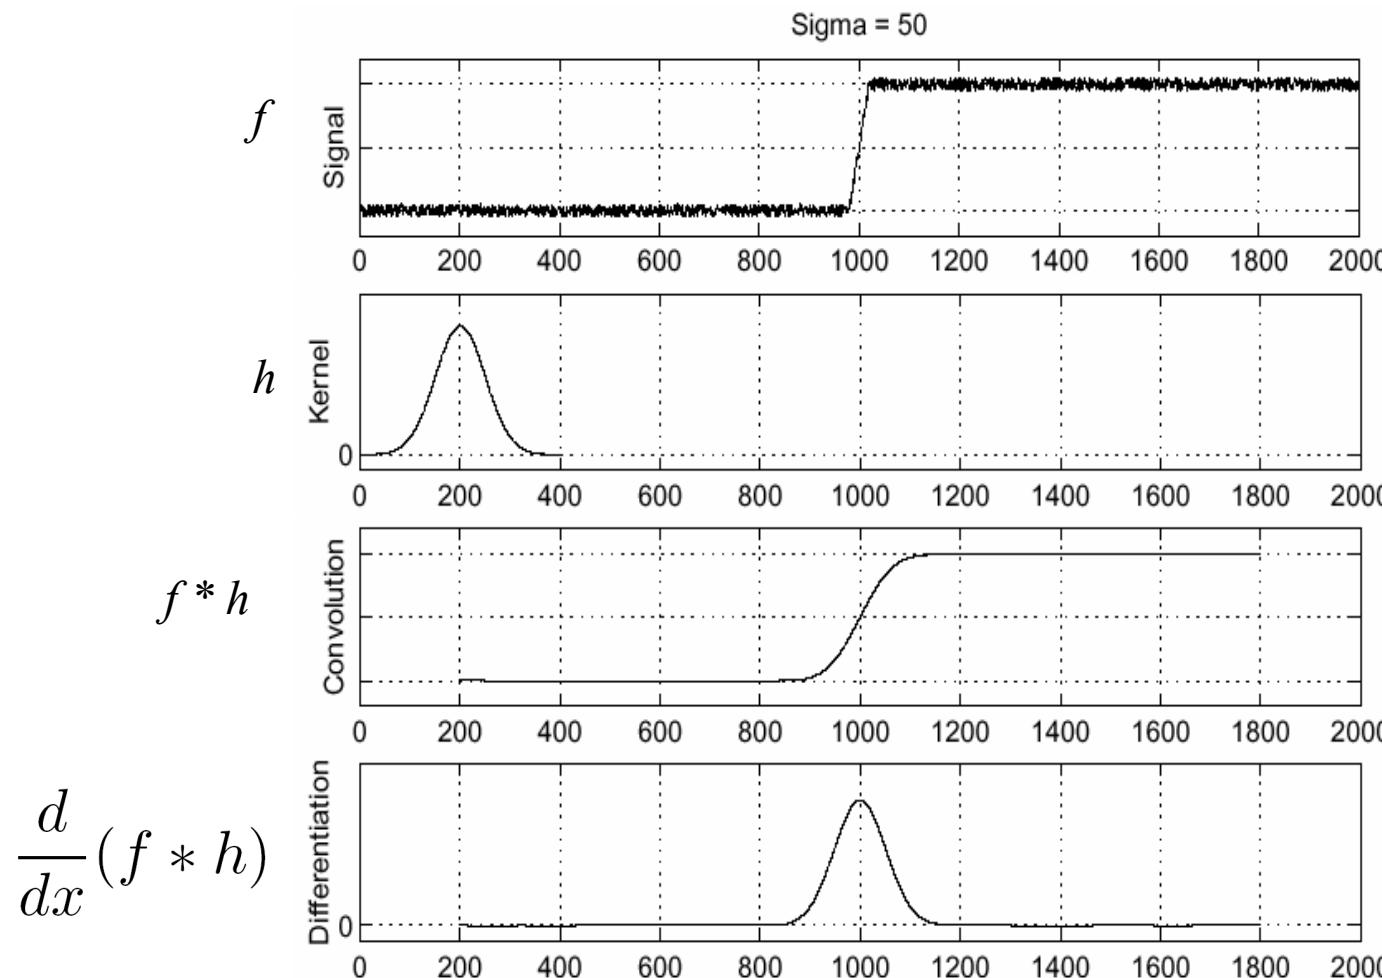


$$\frac{d}{dx}f(x)$$



Where is the edge?

Solution: smooth first



To find edges, look for peaks in $\frac{d}{dx}(f * h)$

Associative property of convolution

- Differentiation is convolution, and convolution is associative:

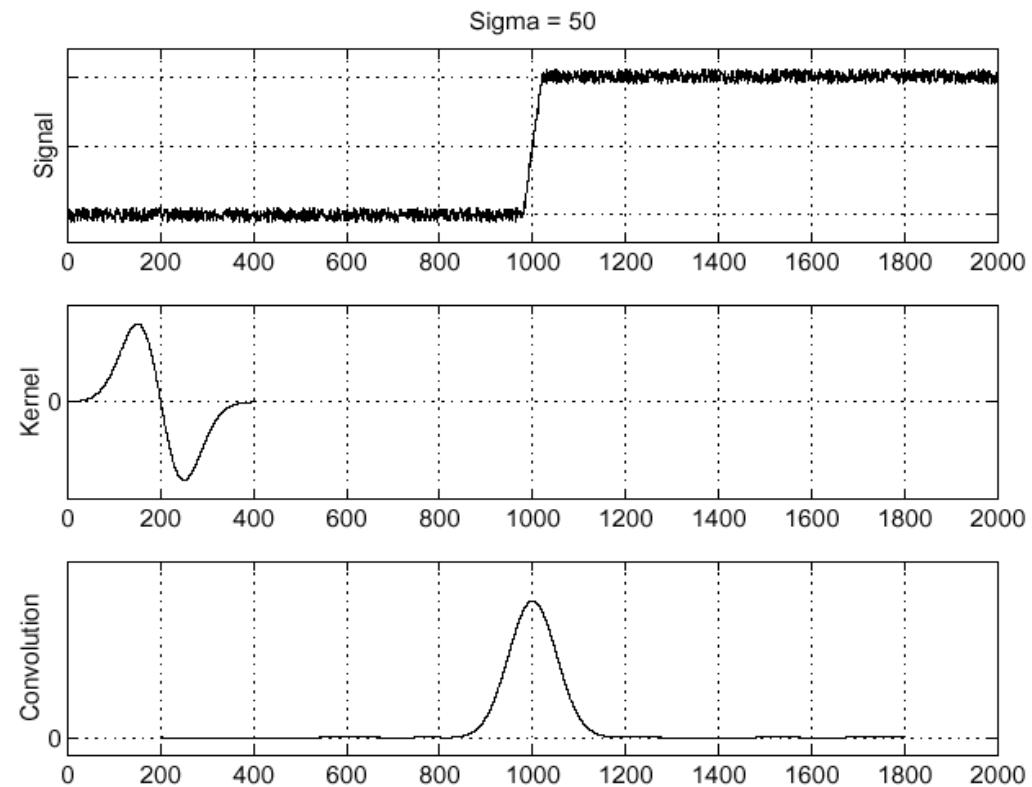
$$\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$$

- This saves us one operation:

f

$$\frac{d}{dx}h$$

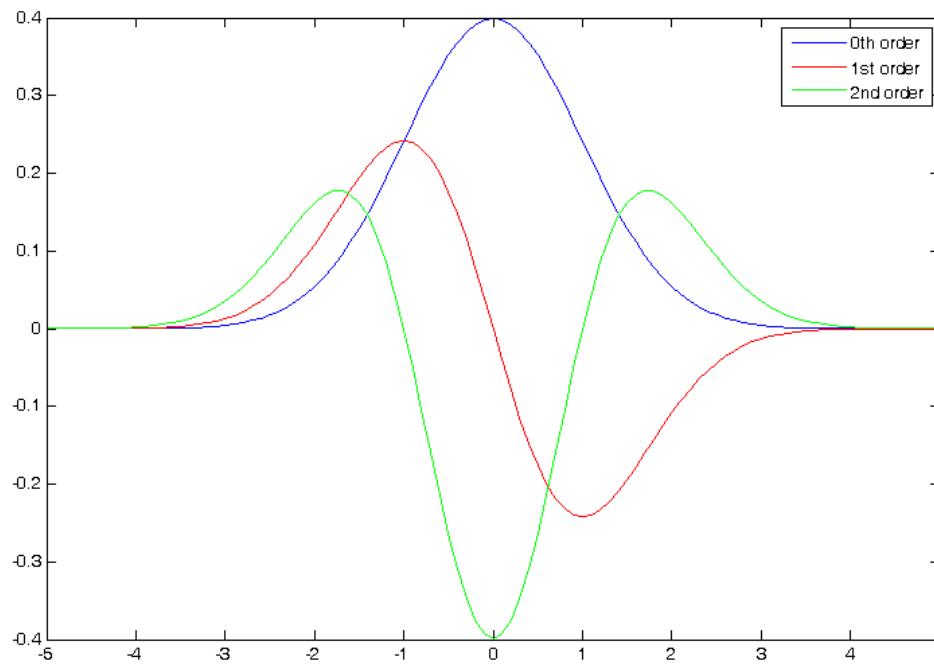
$$f * \frac{d}{dx}h$$



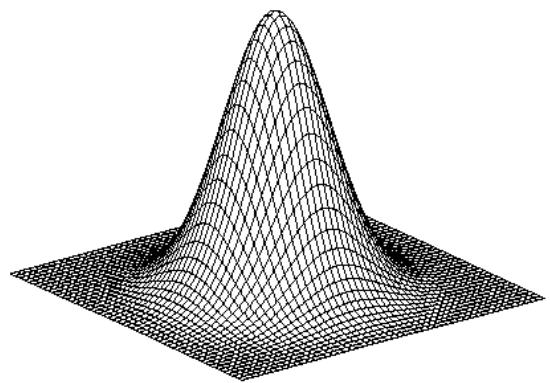
The 1D Gaussian and its derivatives

$$G_\sigma(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

$$G'_\sigma(x) = \frac{d}{dx} G_\sigma(x) = -\frac{1}{\sigma} \left(\frac{x}{\sigma}\right) G_\sigma(x)$$

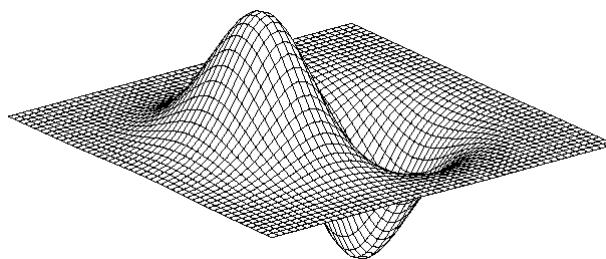


2D edge detection filters



Gaussian

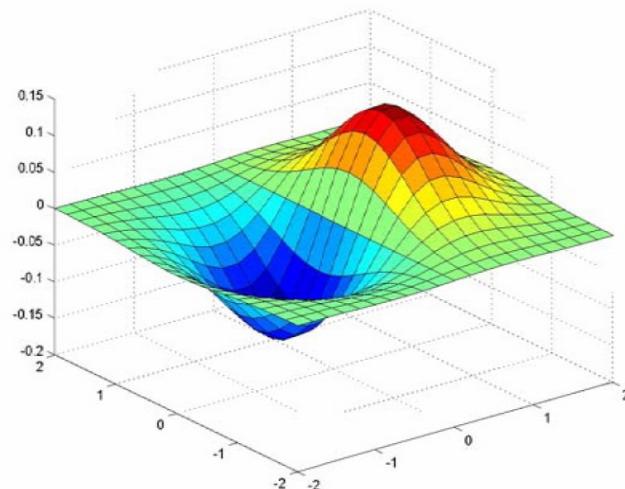
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



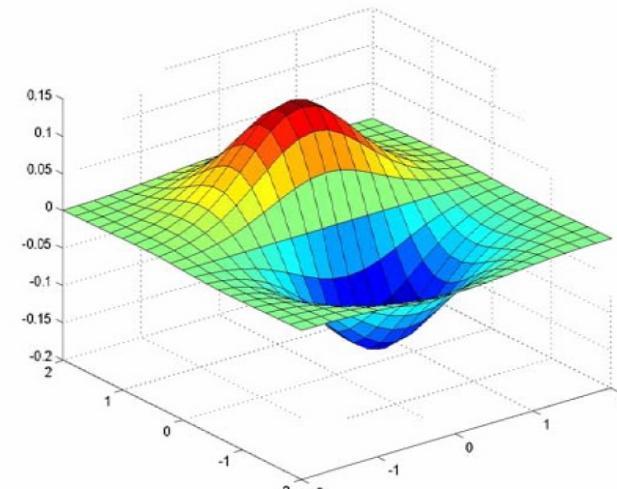
derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

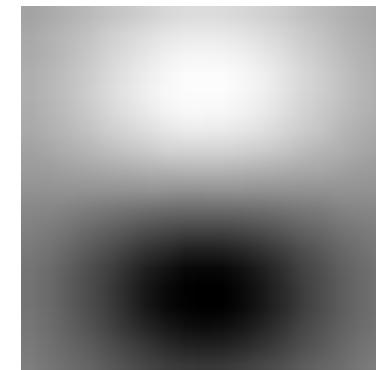
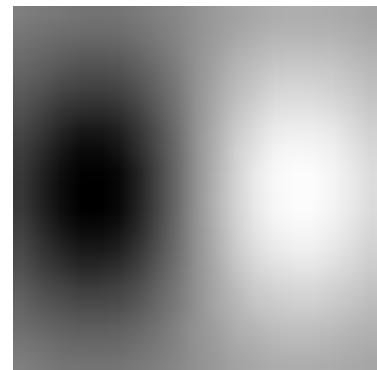
Derivative of Gaussian filter



x-direction



y-direction



The Sobel operator

- Common approximation of derivative of Gaussian

$$\frac{1}{8} \begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$$

s_x

$$\frac{1}{8} \begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix}$$

s_y

- The standard definition of the Sobel operator omits the $1/8$ term
 - doesn't make a difference for edge detection
 - the $1/8$ term **is** needed to get the right gradient magnitude

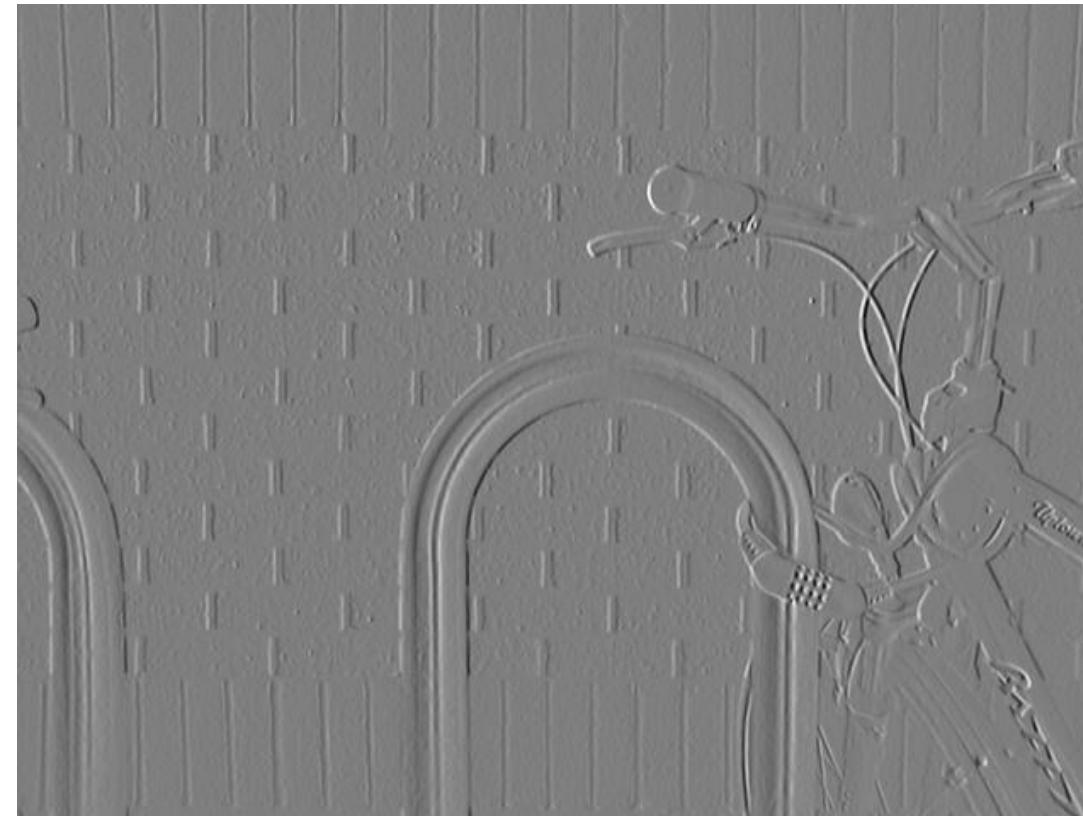
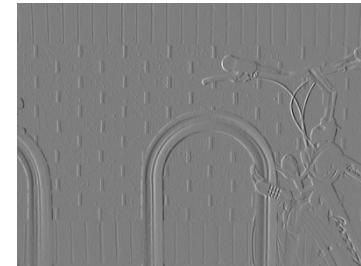
Sobel operator: example

https://en.wikipedia.org/wiki/Sobel_operator



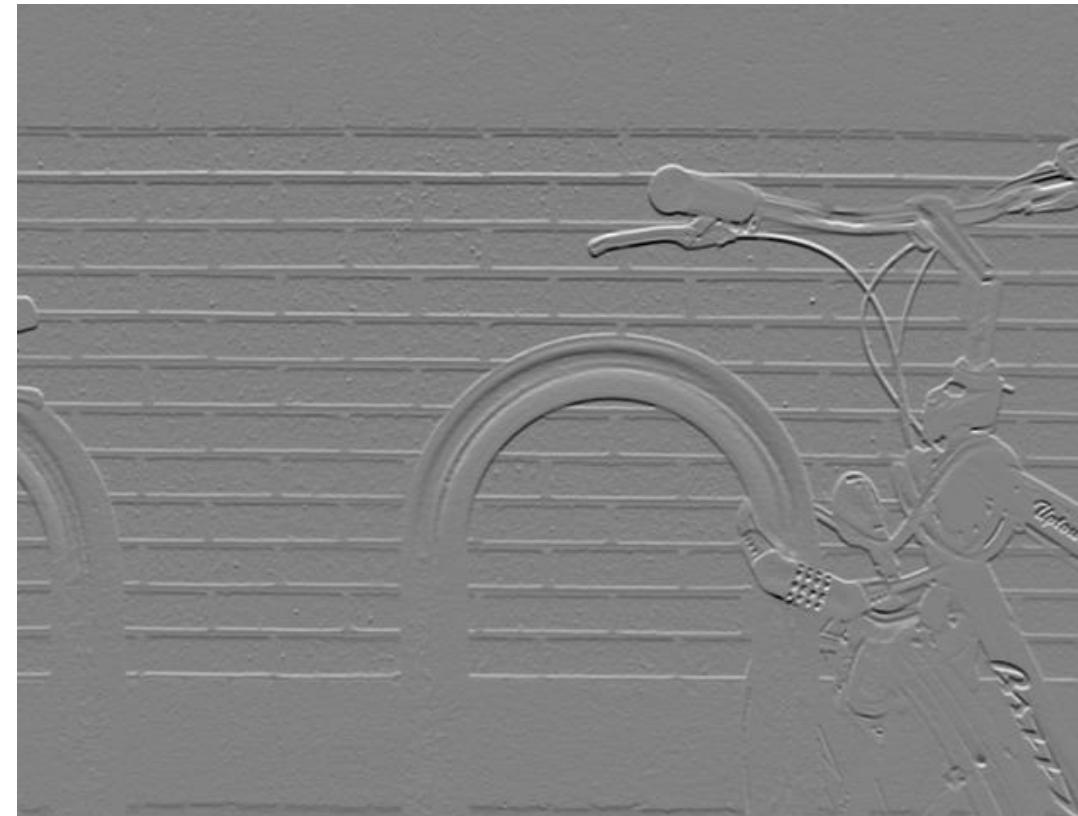
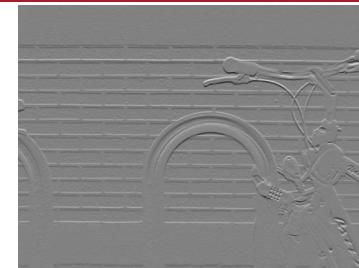
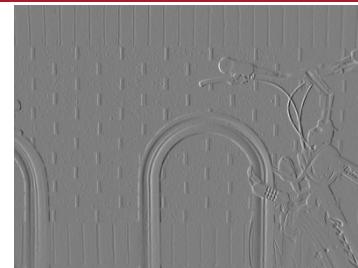
Sobel operator: example

https://en.wikipedia.org/wiki/Sobel_operator



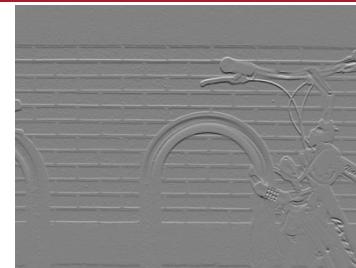
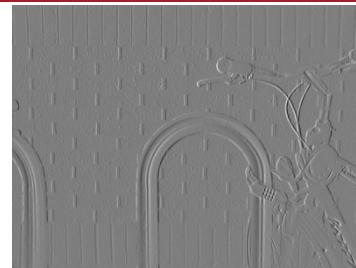
Sobel operator: example

https://en.wikipedia.org/wiki/Sobel_operator



Sobel operator: example

https://en.wikipedia.org/wiki/Sobel_operator



Canny Edge detection



original image

Finding edges



smoothed gradient magnitude

Finding edges



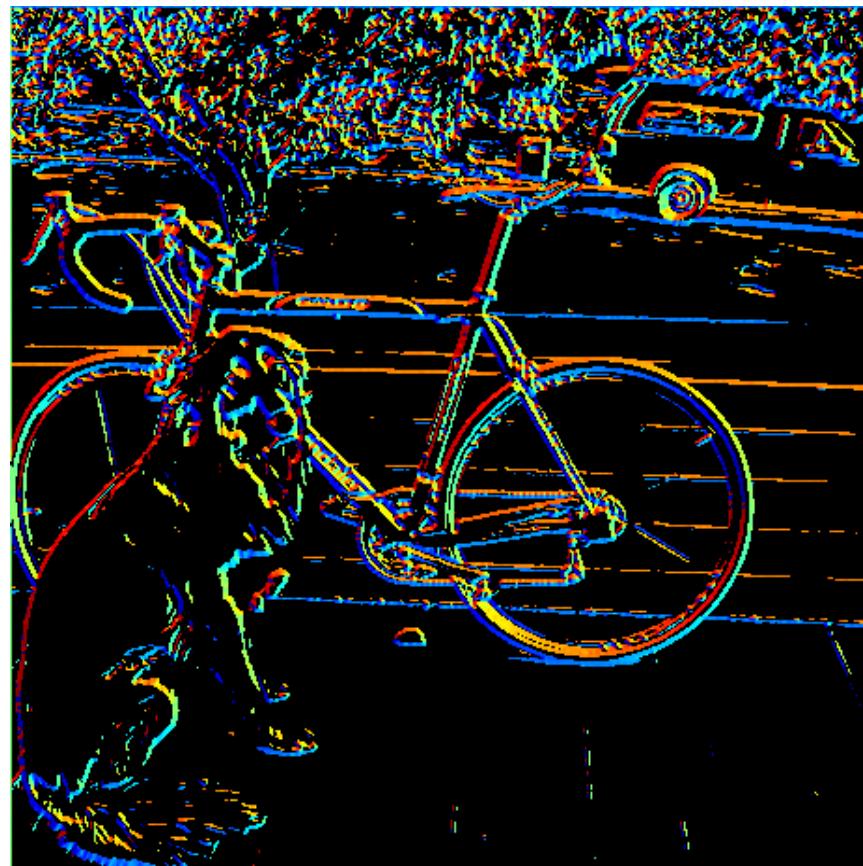
where is the edge?

smoothed gradient magnitude

Get Orientation at Each Pixel

- Get orientation (below, threshold at minimum gradient magnitude)

$$\theta = \text{atan2}(g_y, g_x)$$

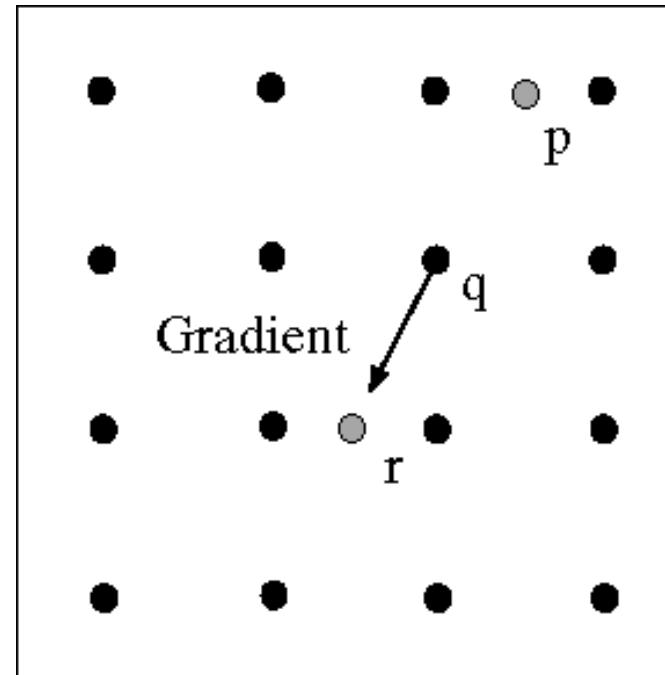
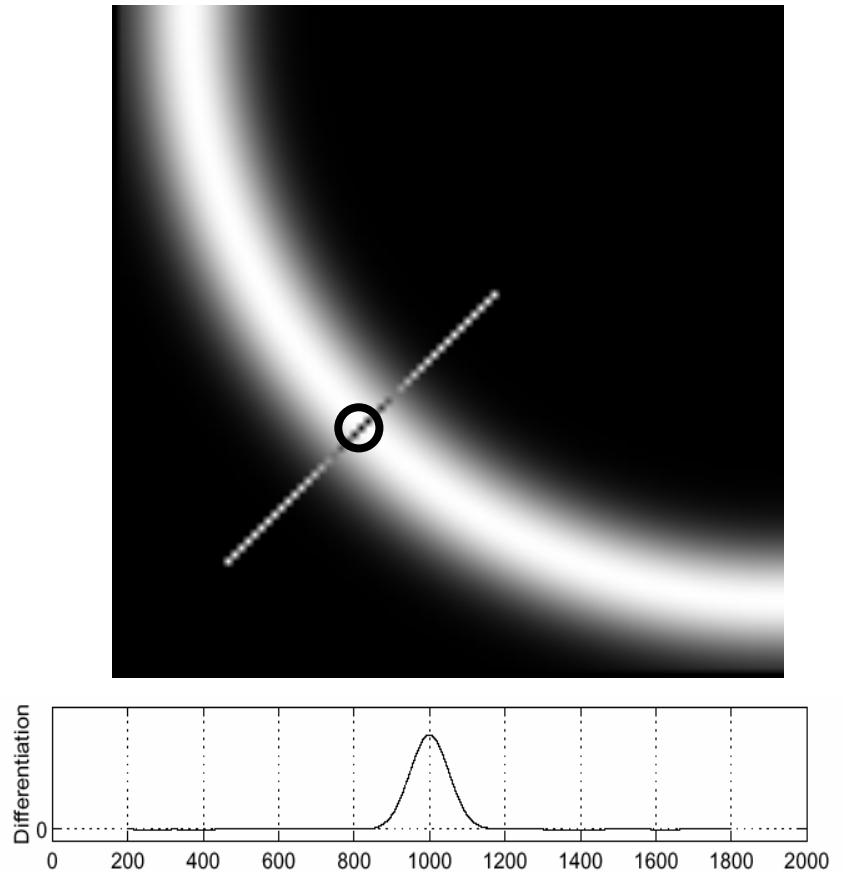


360

Gradient orientation angle

0

Non-maximum suppression



- Check if pixel is local maximum along gradient direction
 - requires *interpolating* pixels p and r

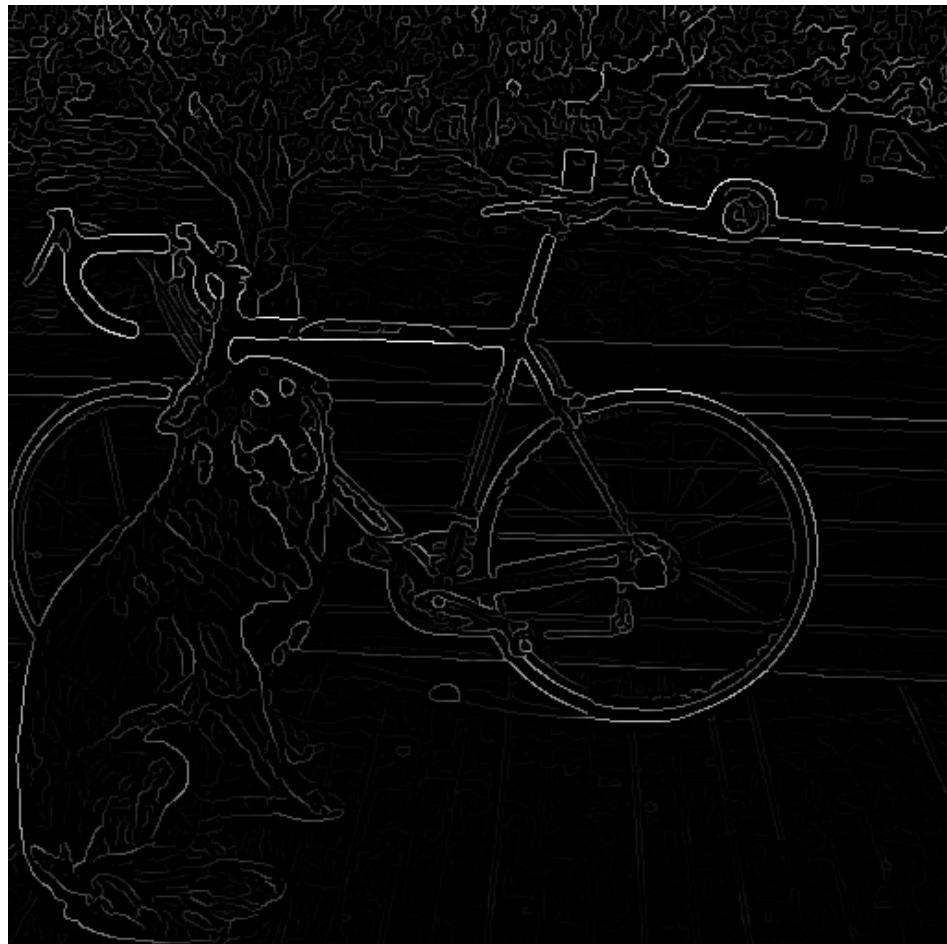
Non-maximum suppression - Before



Non-maximum suppression - After

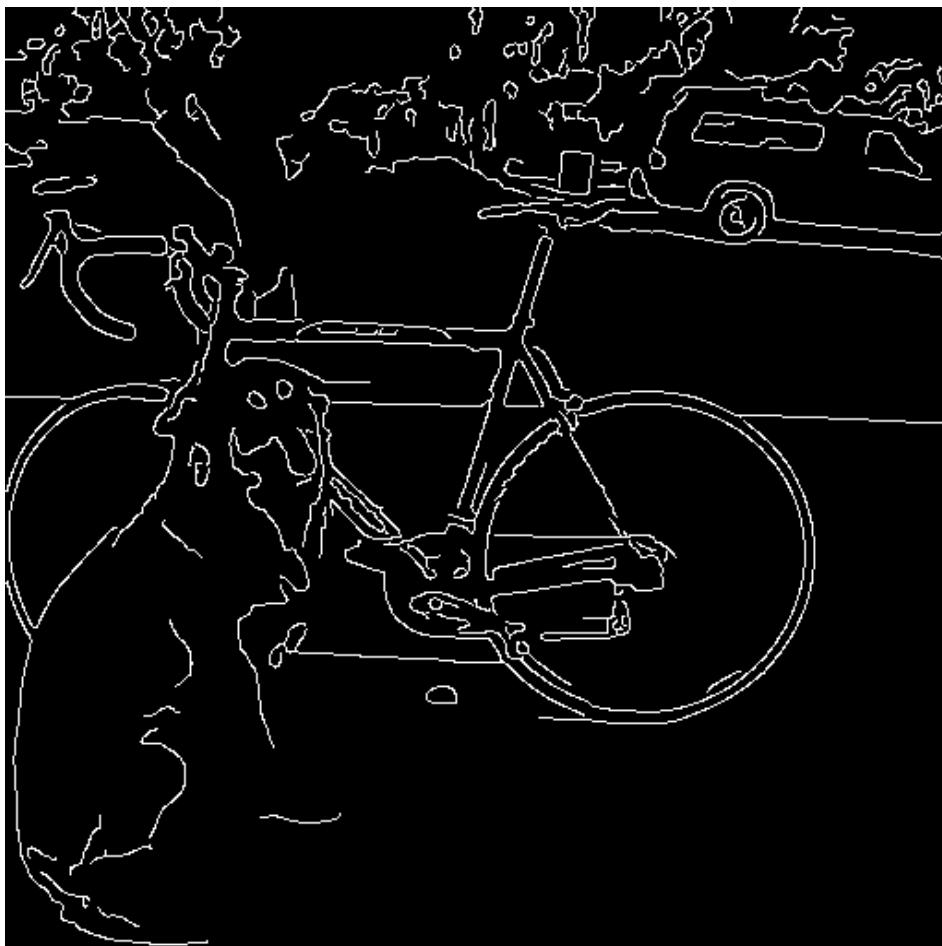


Thresholding edges



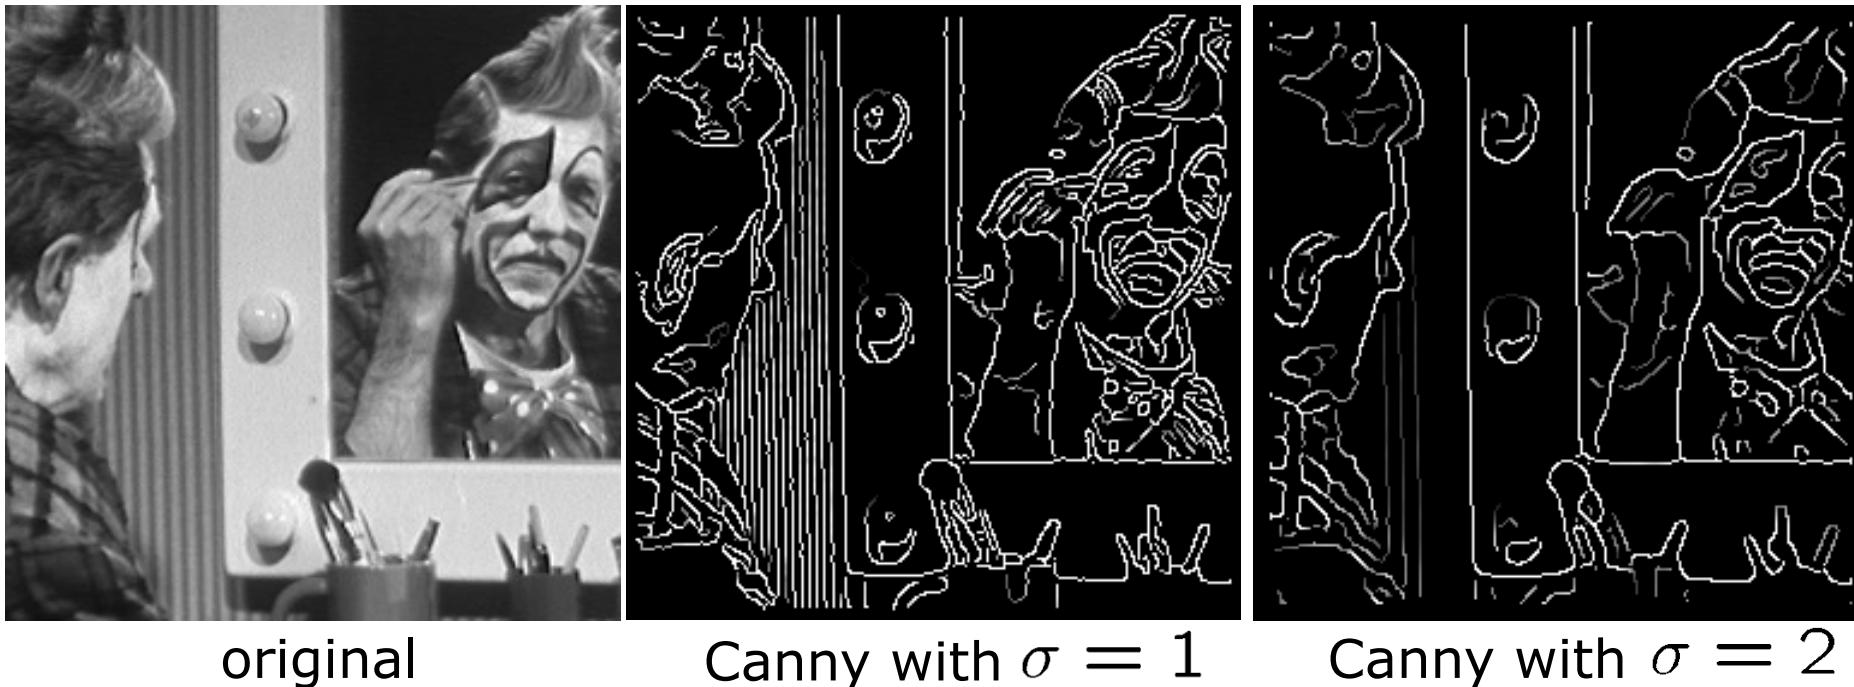
- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
 - $R > T$: strong edge
 - $R < T$ but $R > t$: weak edge
 - $R < t$: no edge
- Why two thresholds?

Connecting edges



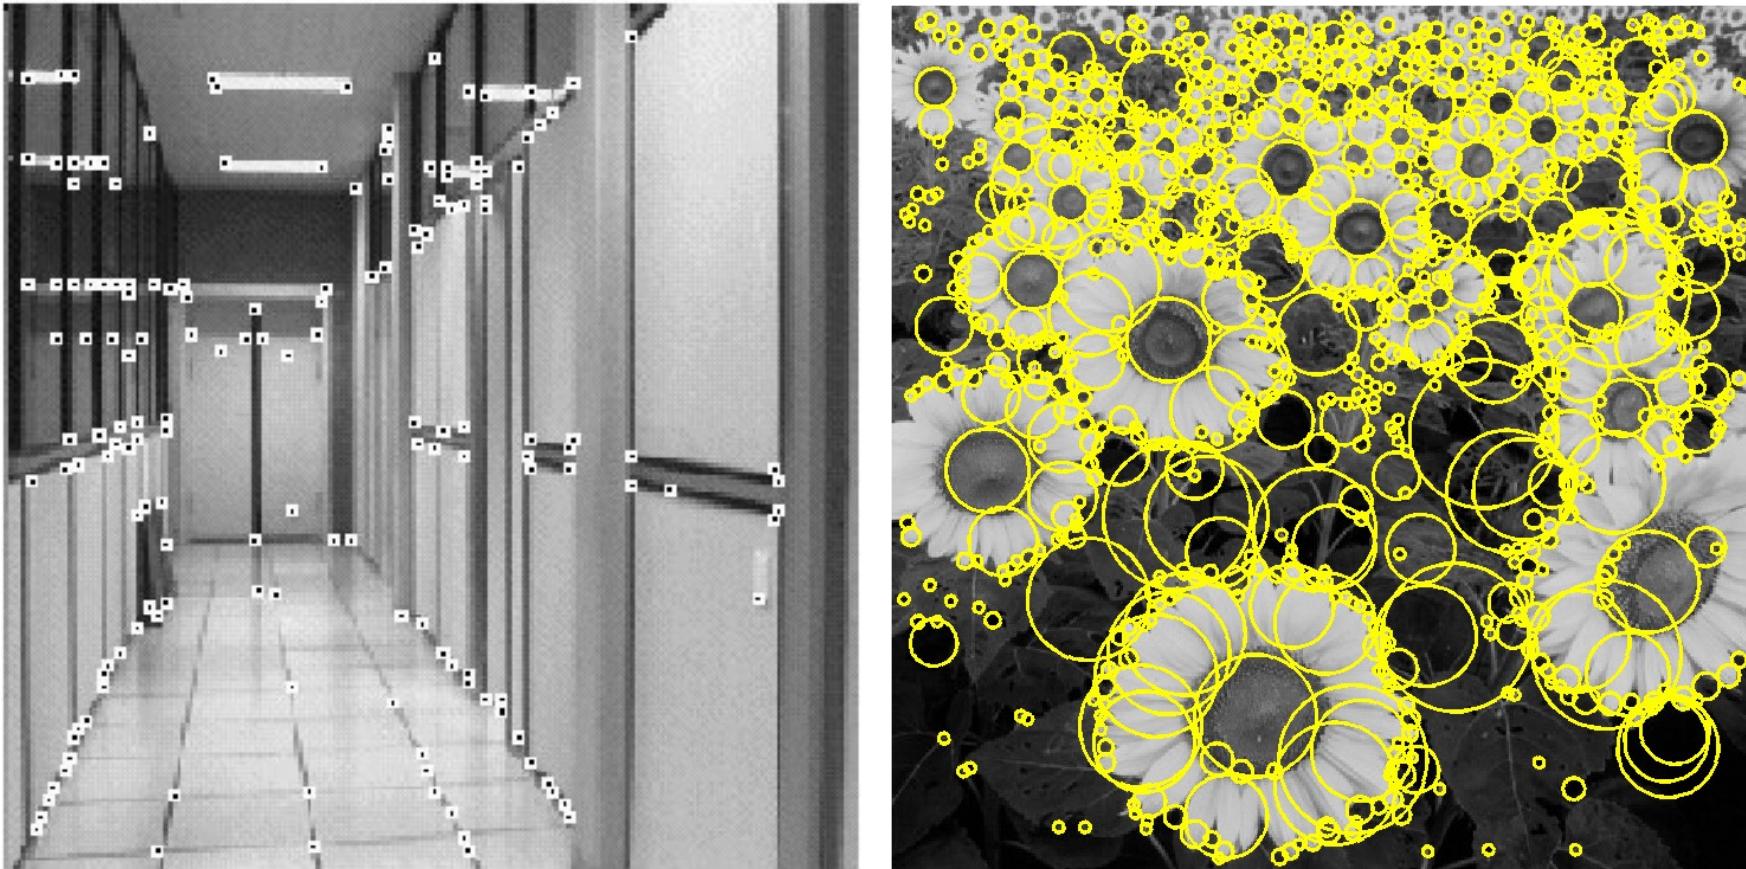
- Strong edges are edges!
- Weak edges are edges if they connect to strong
- Look in some neighborhood (usually 8 closest)

Canny edge detector



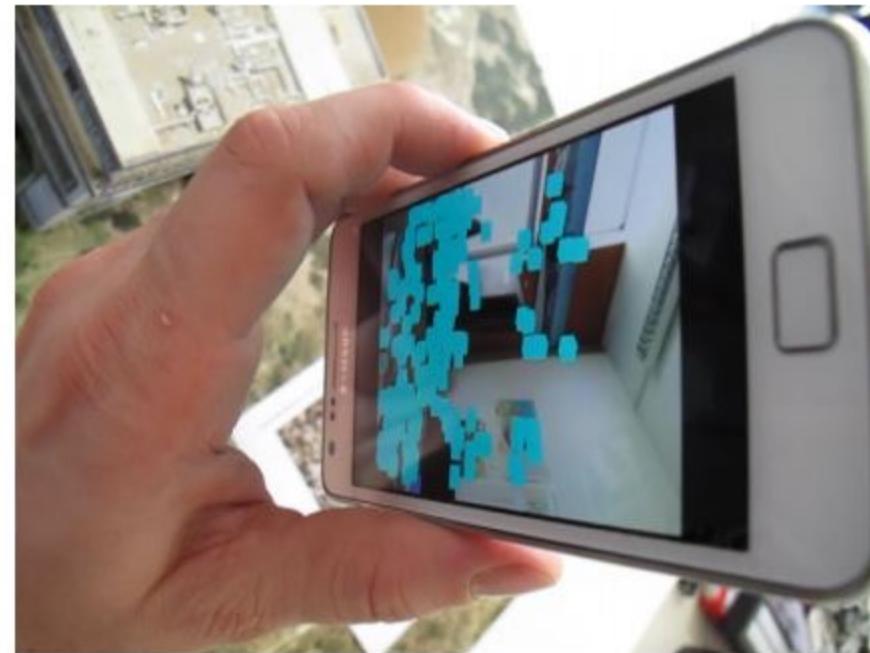
- The choice of σ depends on desired behavior
 - large σ detects “large-scale” edges
 - small σ detects fine edges

Feature extraction - Corners and blobs

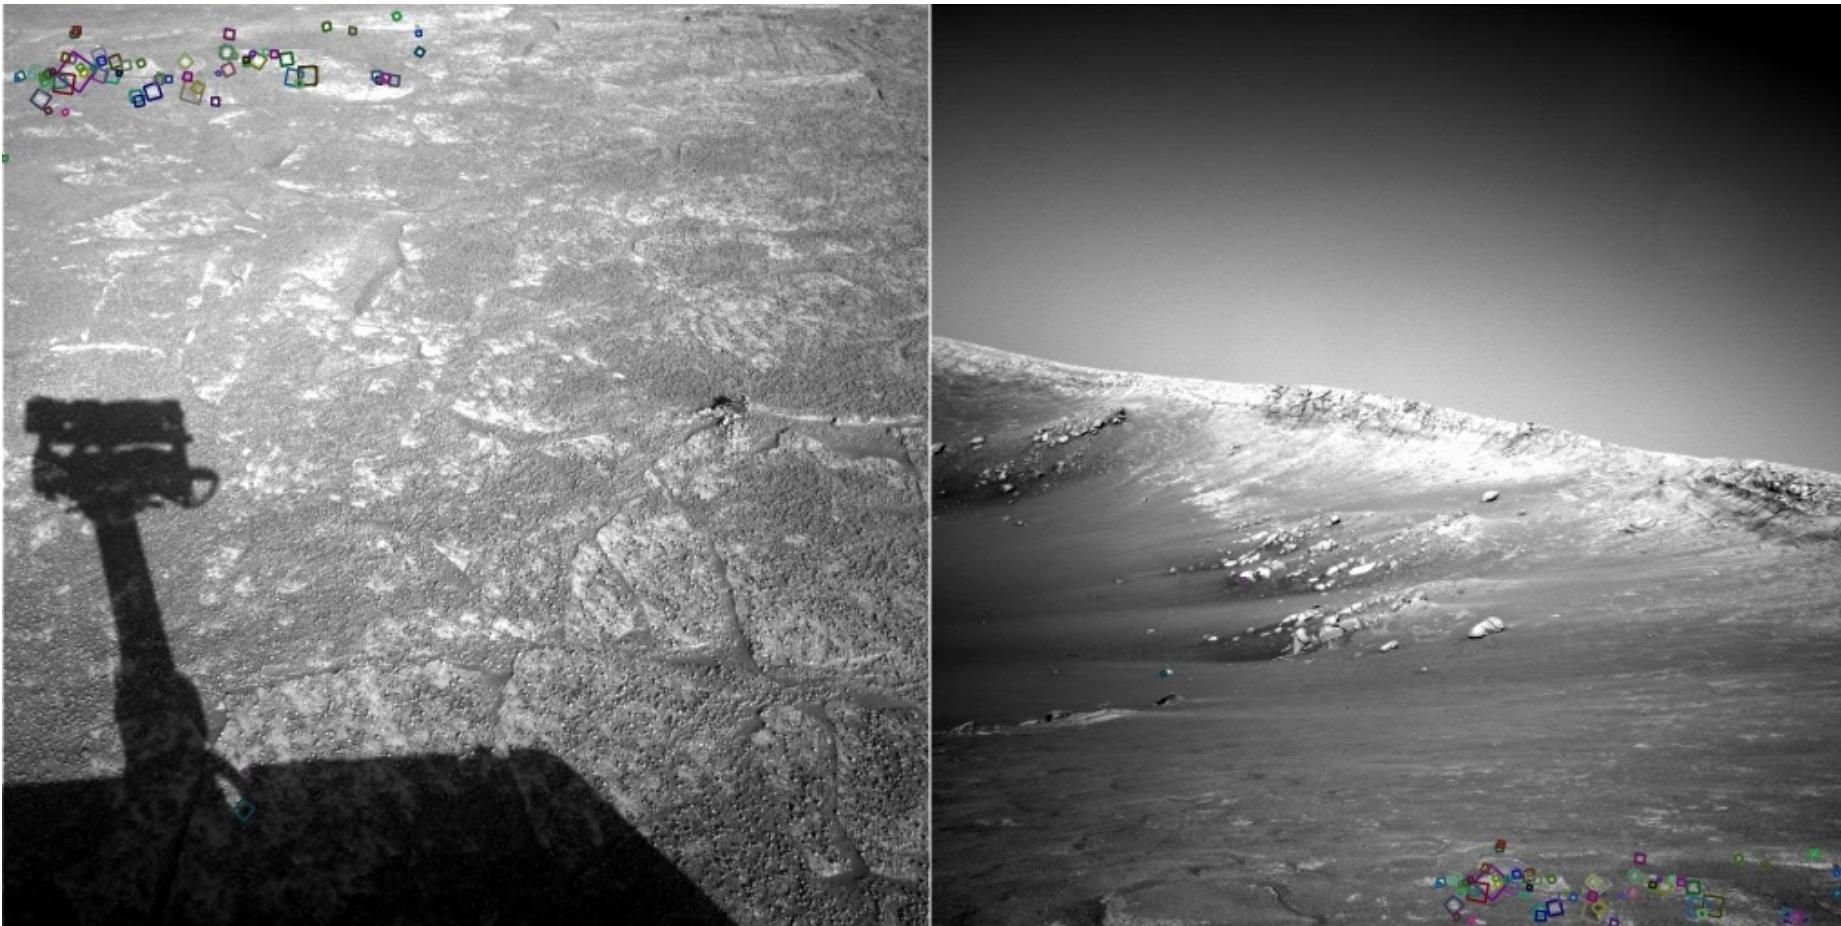


Application - Visual SLAM

- (aka Simultaneous Localization and Mapping)

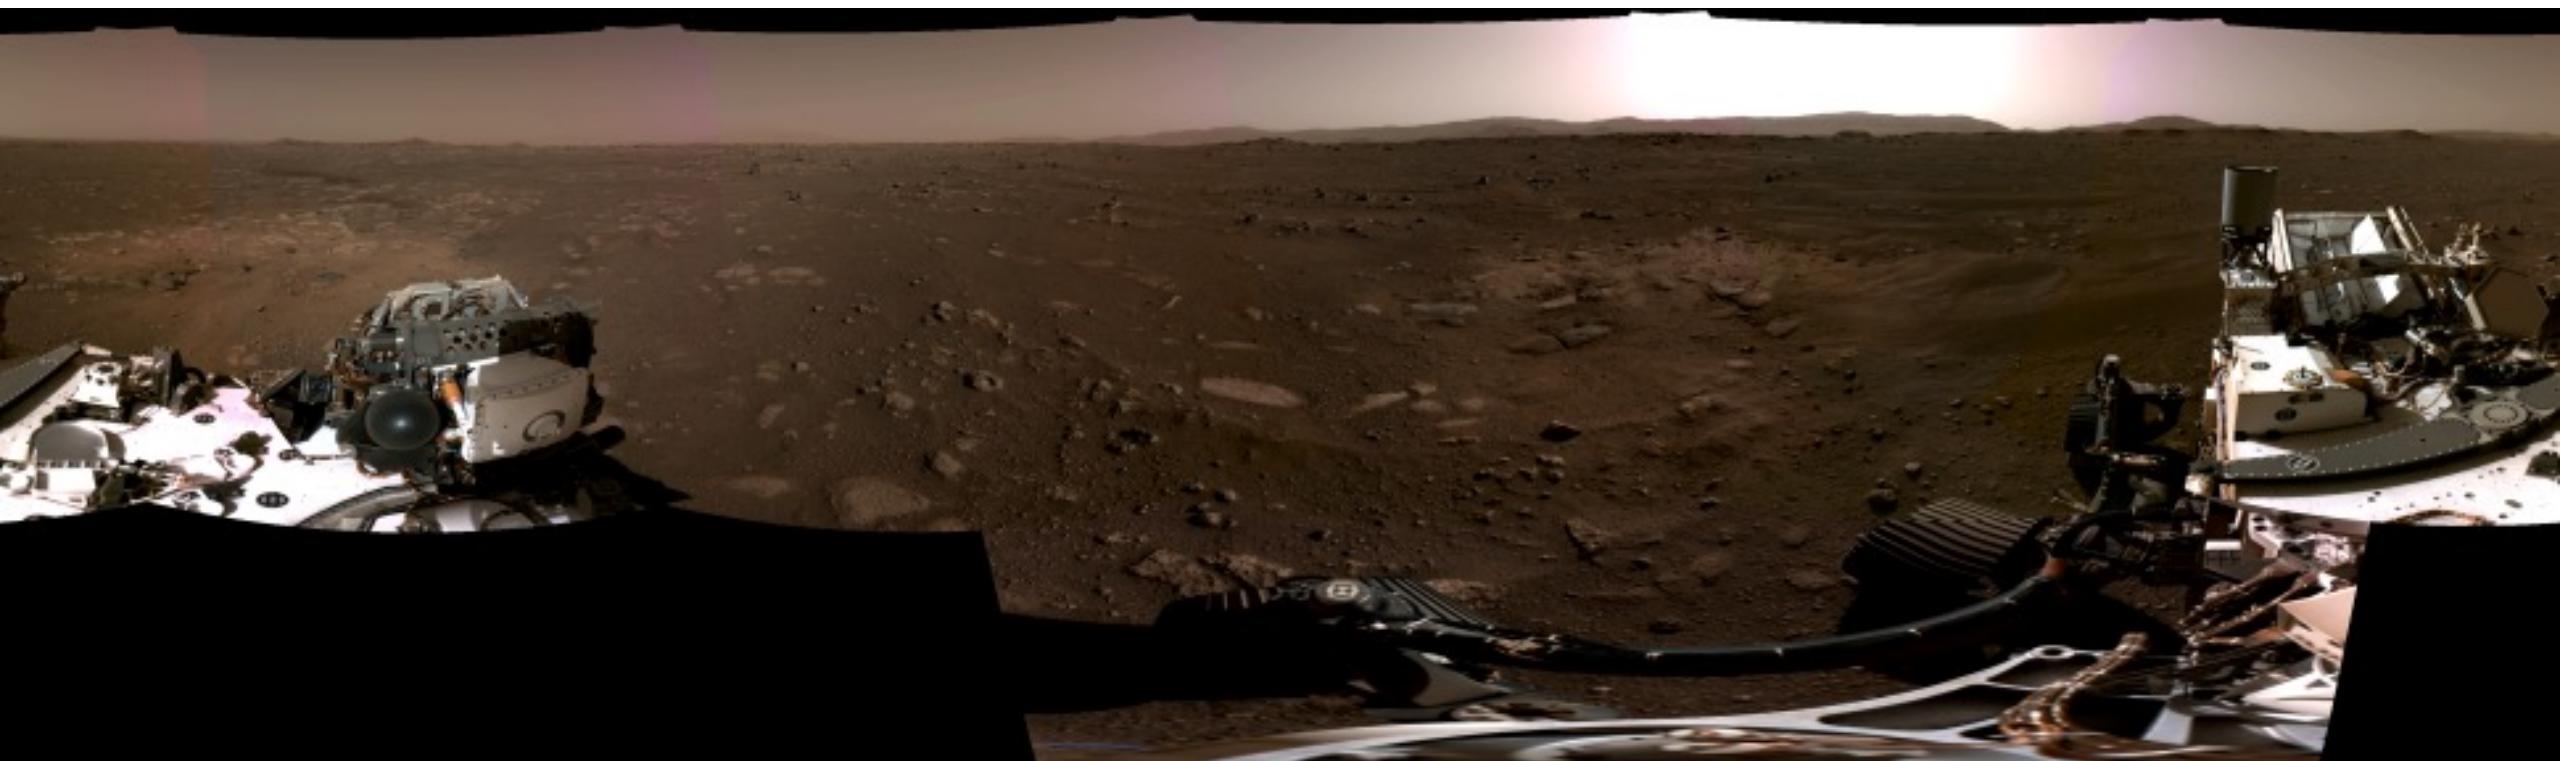


Application - Mars exploration



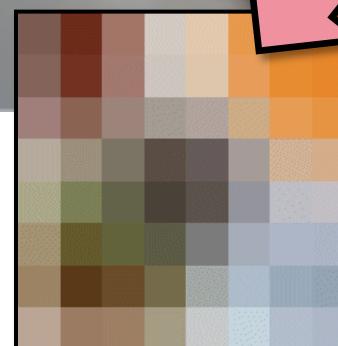
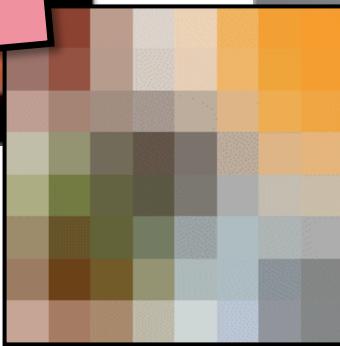
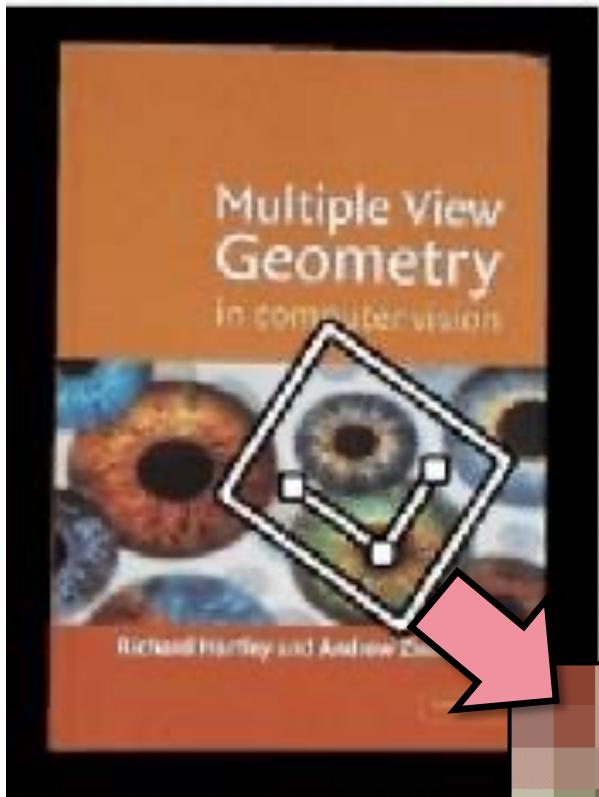
NASA Mars Rover images
with SIFT feature matches

Application - Panorama stitching



Panorama captured by Perseverance Rover, Feb. 20, 2021

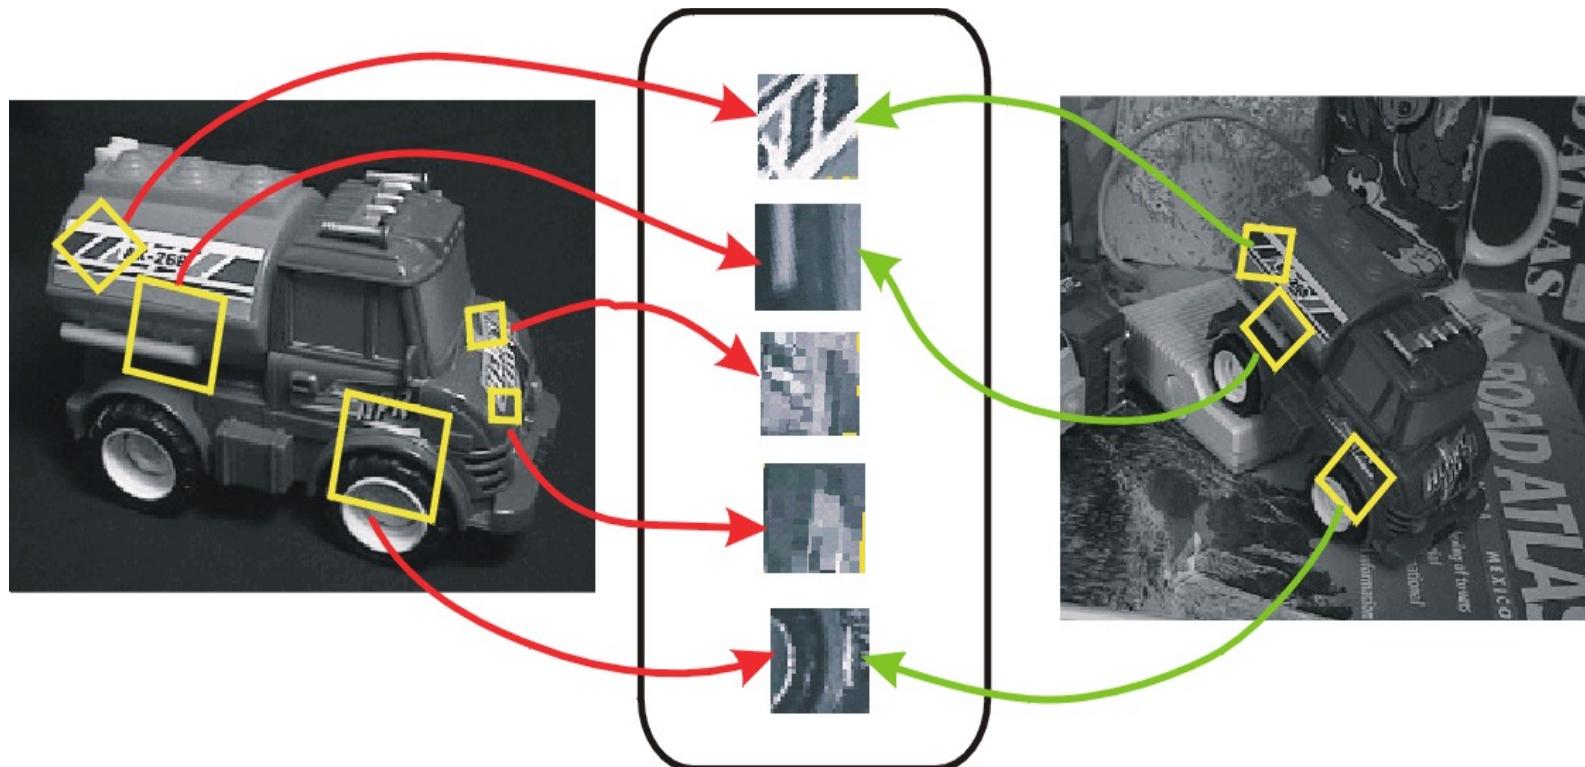
Application - Object search



Invariant local features

Find features that are invariant to transformations

- geometric invariance: translation, rotation, scale
- photometric invariance: brightness, exposure, ...



Feature Descriptors

More motivation ...

Feature points are used for:

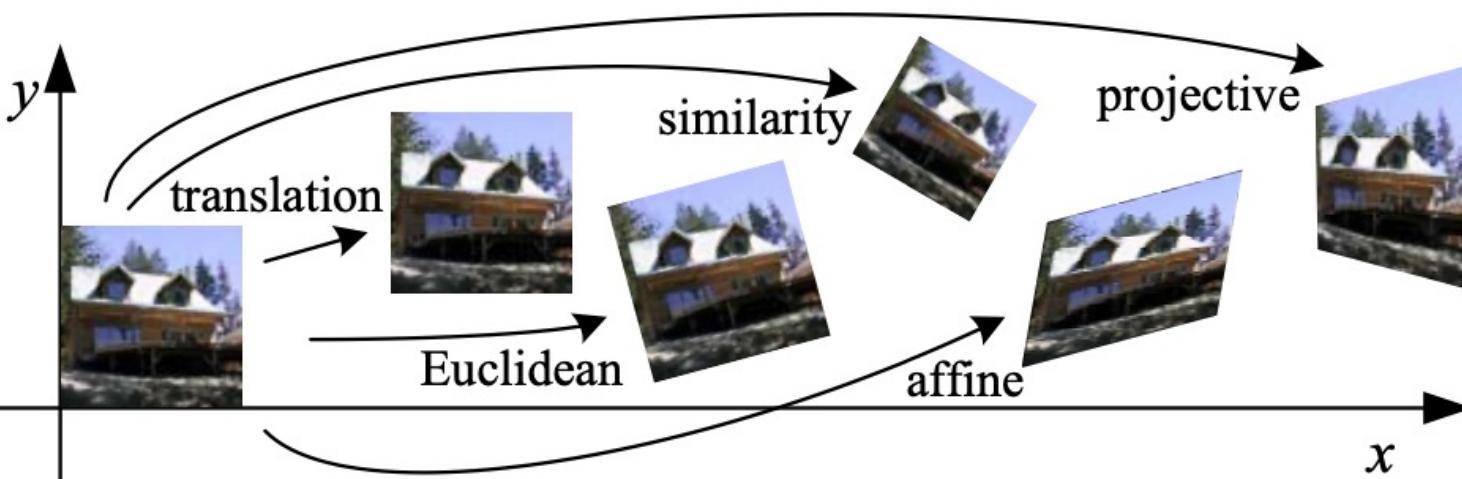
- Image alignment (e.g., mosaics)
- 3D reconstruction
- Motion tracking (e.g., for AR)
- Object recognition
- Image retrieval
- Robot/car navigation



Feature extraction algorithms

- Canny Edge Detector
- Harris Corner Detector
- Shi-Tomasi Corner Detector
- SIFT (Scale-Invariant Feature Transform)
- SURF (Speed-Up Robust Features)
- ORB (Oriented FAST and Rotated BRIEF)
- HOG (Histogram of Oriented Gradients)
- LBP (Local Binary Pattern)
- ORB-SLAM
- ...

Image transformations



Transformation	Matrix	# DoF	Preserves	Icon
translation	$[\mathbf{I} \quad \mathbf{t}]_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$[\mathbf{R} \quad \mathbf{t}]_{2 \times 3}$	3	lengths	
similarity	$[s\mathbf{R} \quad \mathbf{t}]_{2 \times 3}$	4	angles	
affine	$[\mathbf{A}]_{2 \times 3}$	6	parallelism	
projective	$[\tilde{\mathbf{H}}]_{3 \times 3}$	8	straight lines	

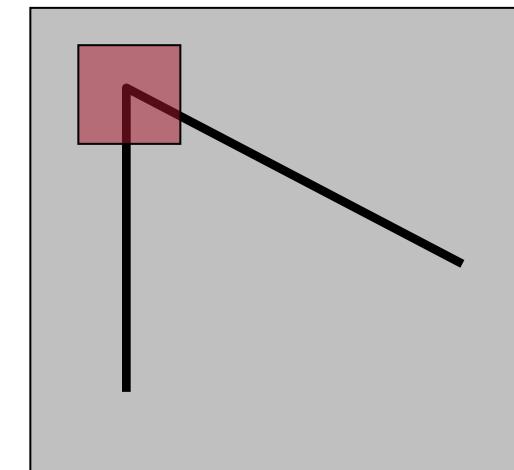
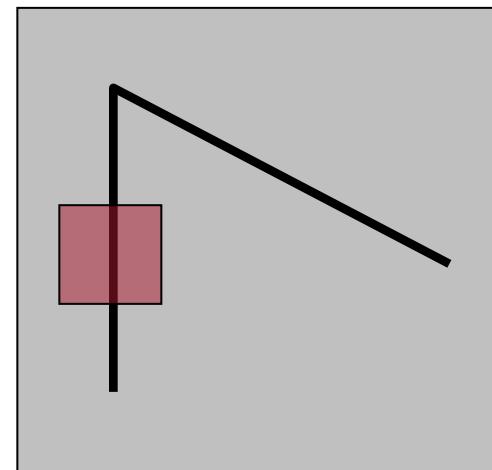
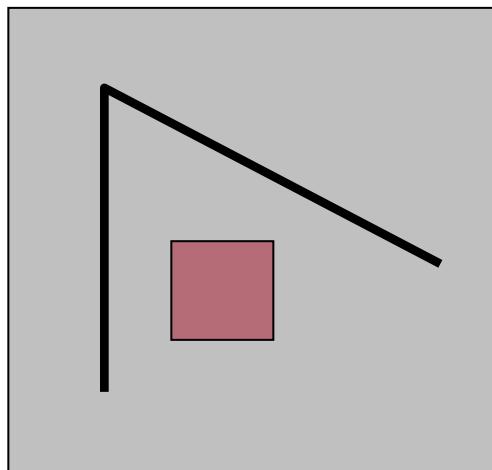
Feature extraction - Corners



Local measures of uniqueness

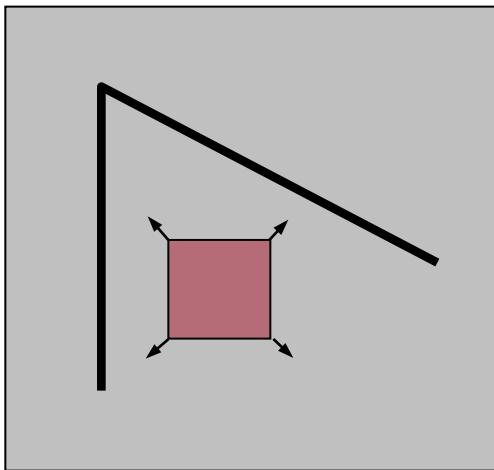
Suppose we only consider a small window of pixels

- What defines whether a feature is a good or bad candidate?

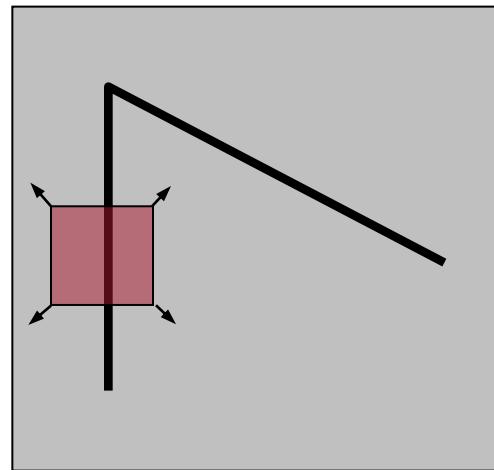


Corner

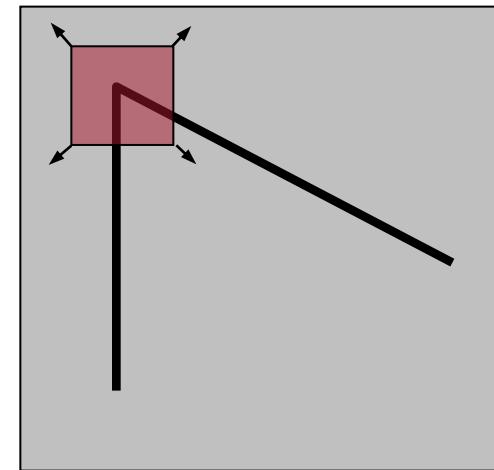
- How does the window change when you shift it?
- Shifting the window in any direction causes a big change



“flat” region:
no change in
all directions



“edge”:
no change along
the edge
direction

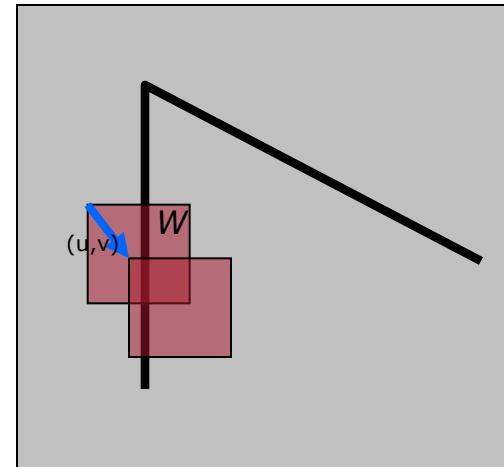


“corner”:
significant
change in all
directions

Corner Detection

- Consider shifting the window W by (u, v)
- Compare each pixel before and after by summing up the squared differences (SSD)
- This defines an SSD “error” $E(u, v)$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$



Taylor Series / Expansion - Reminder

Taylor Series expansion of I :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion (u, v) is small, then first order approximation is good

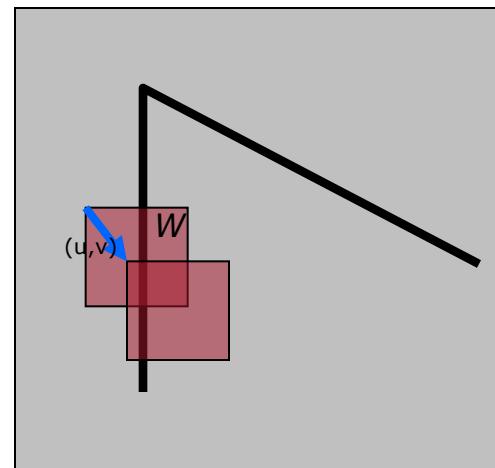
$$\begin{aligned} I(x + u, y + v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \\ &\qquad\qquad\qquad \text{shorthand: } I_x = \frac{\partial I}{\partial x} \end{aligned}$$

Plugging this into the formula on the previous slide...

Corner Detection

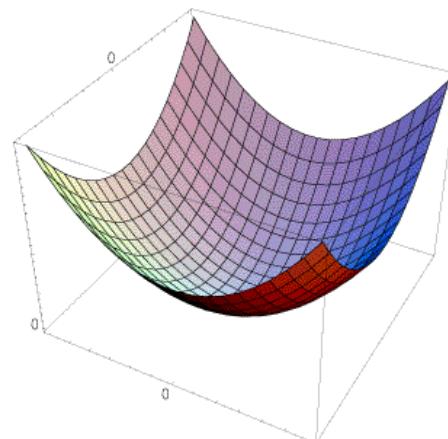
- Consider shifting the window W by (u, v)
- Compare each pixel before and after by summing up the squared differences (SSD)
- This defines an SSD “error” $E(u, v)$:

$$E(u, v) = \sum_{x,y} w(x, y)[I(x+u, y+v) - I(x, y)]^2$$



- **Apply the Taylor Expansion:**

$$E(u, v) \approx [u \quad v] \left(\sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$



Corner Detection

$$M = \sum w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

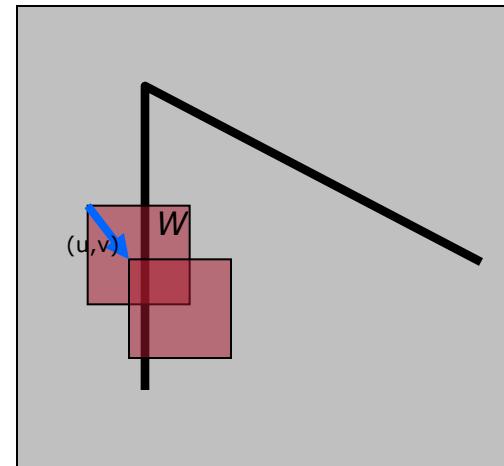
$$R = \text{Det}(M) - k \text{Trace}^2(M)$$

k is an empirically determined constant; $k \in [0.04, 0.06]$

Where

$$\text{Det}(M) = \lambda_1 \lambda_2$$

$$\text{Trace}(M) = \lambda_1 + \lambda_2$$

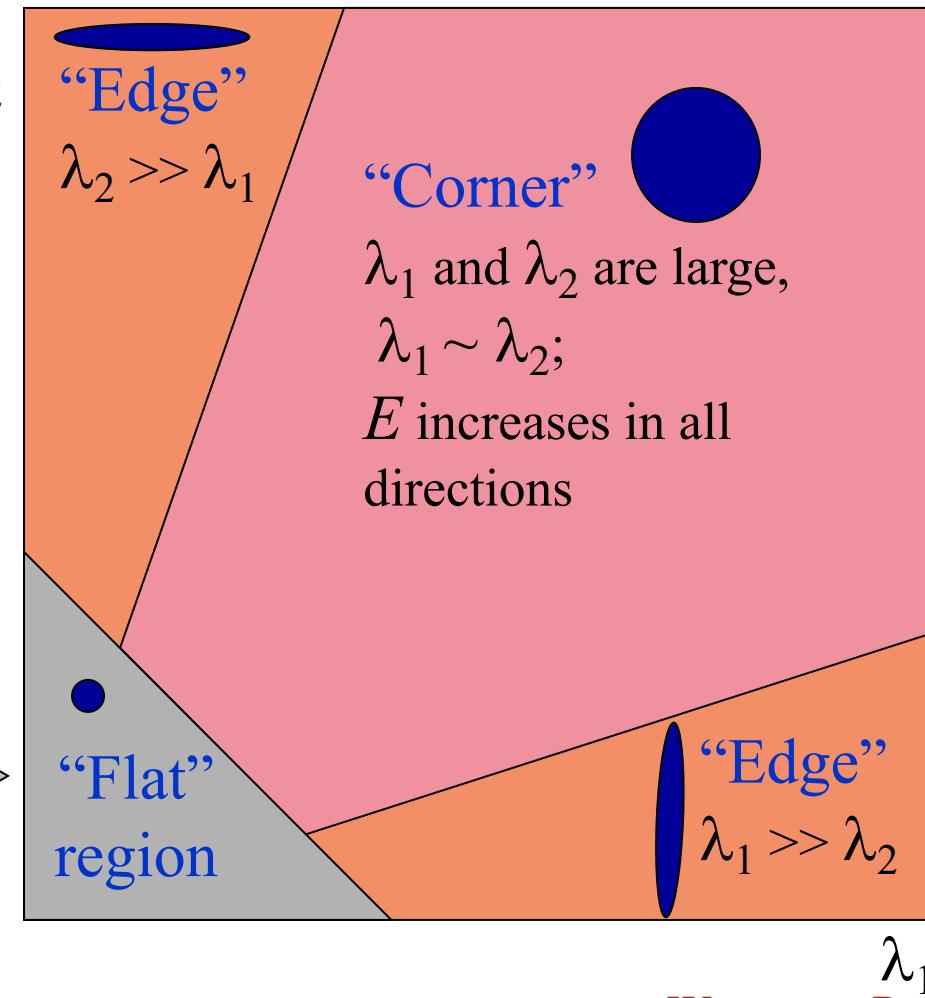


Interpreting the eigenvalues

Classification of image points using eigenvalues of M :

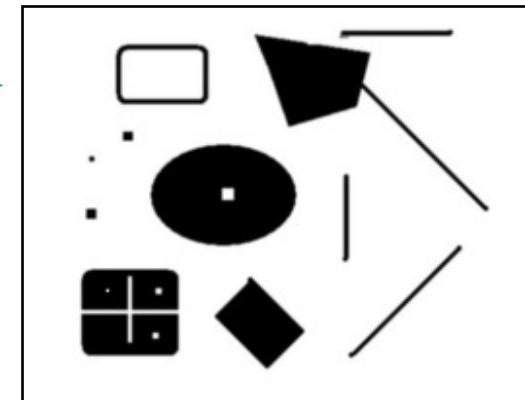
$$E(u,v) \approx [u \quad v] \left(\sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$

λ_1 and λ_2 are small;
 E is almost constant
in all directions

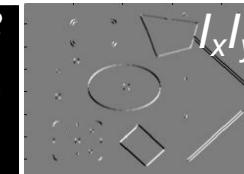
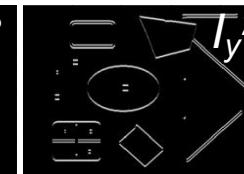
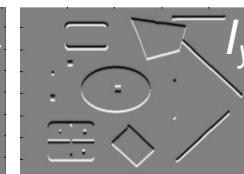
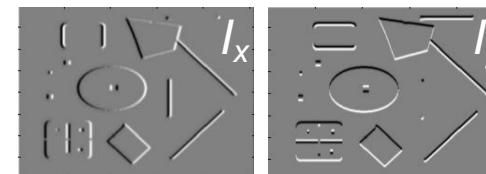


Harris Detector [Harris88] - Steps

0. Input image



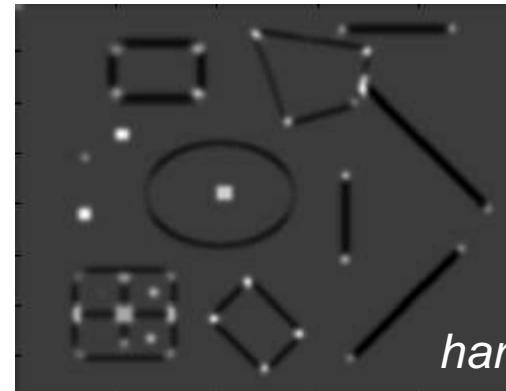
1. Image derivatives



2. Square of derivatives



3. Gaussian filter $G(s_I)$



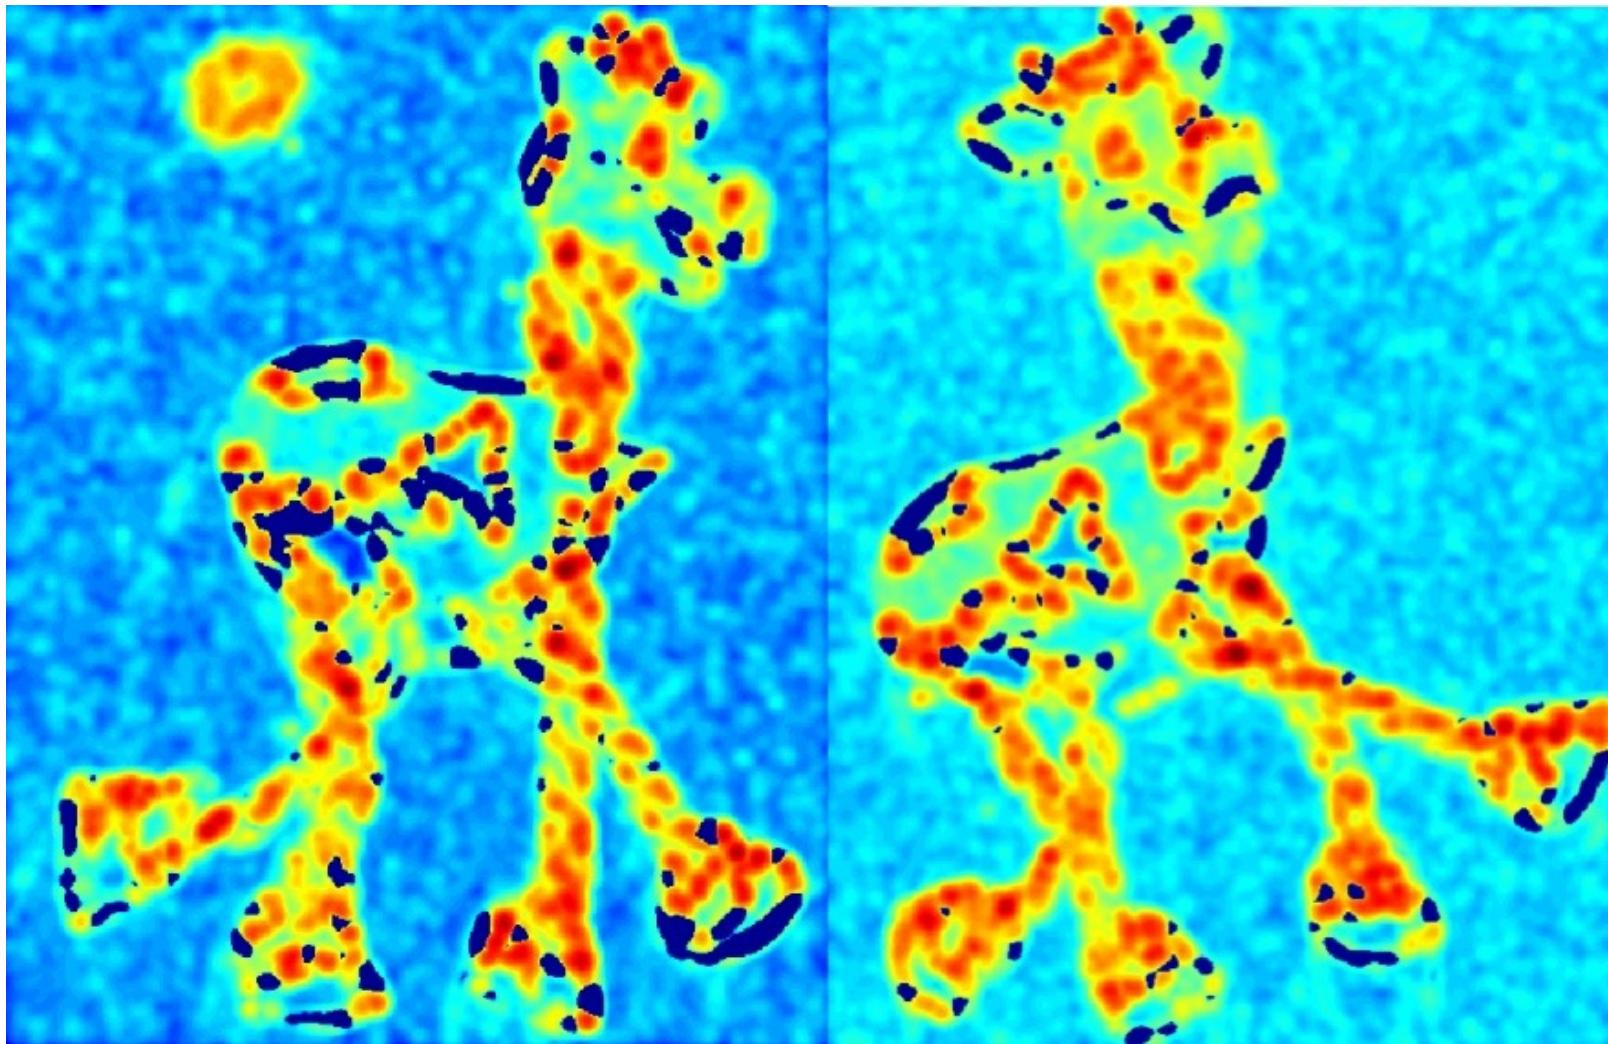
4. Cornerness function - both eigenvalues are strong

5. Non-maxima suppression

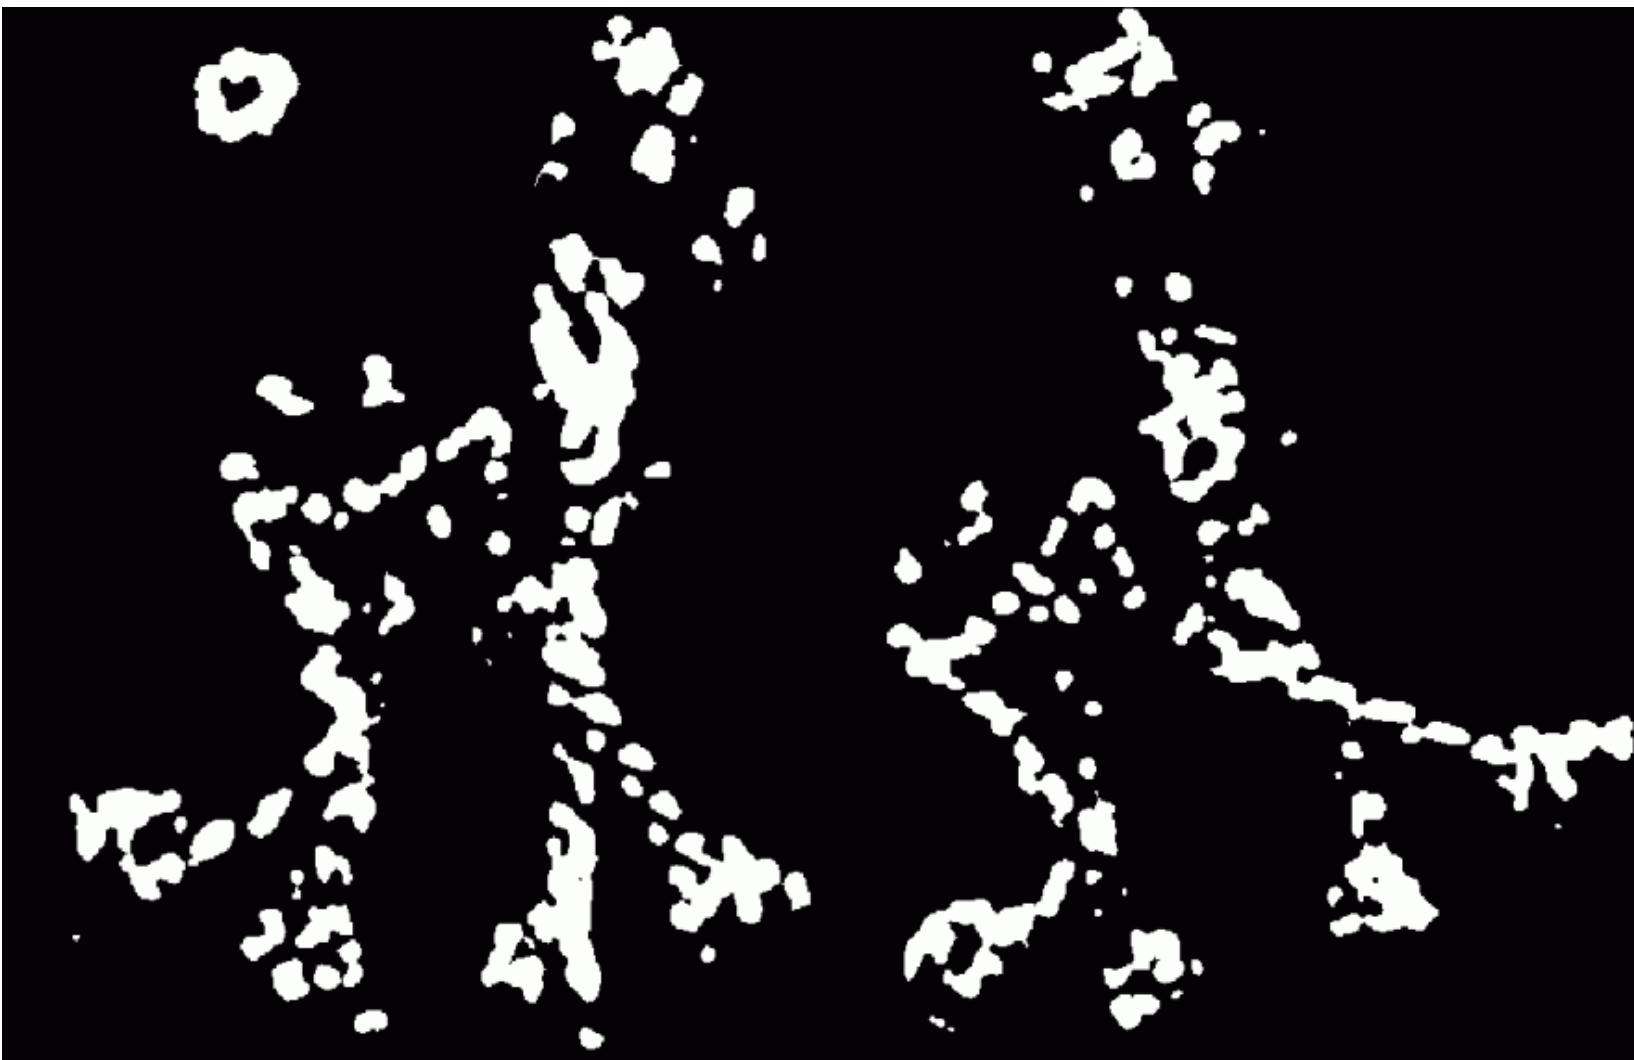
Harris detector example



R value (red high, blue low)



Threshold ($R > \text{value}$)



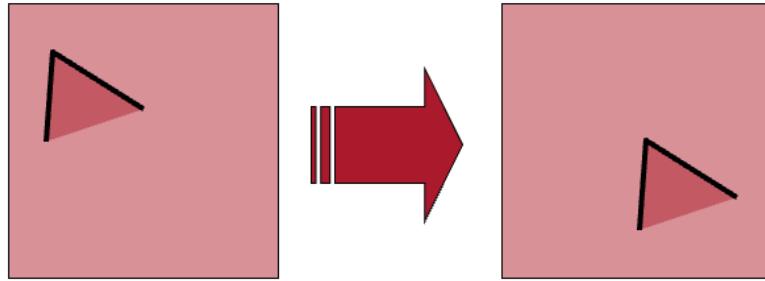
Find local maxima of R (non-max suppression)



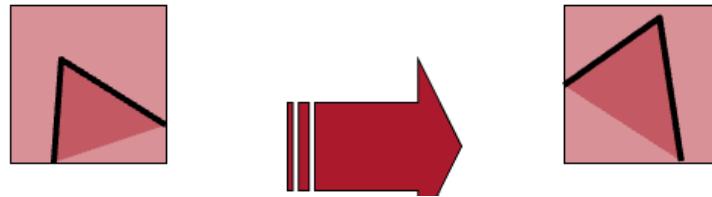
Harris features



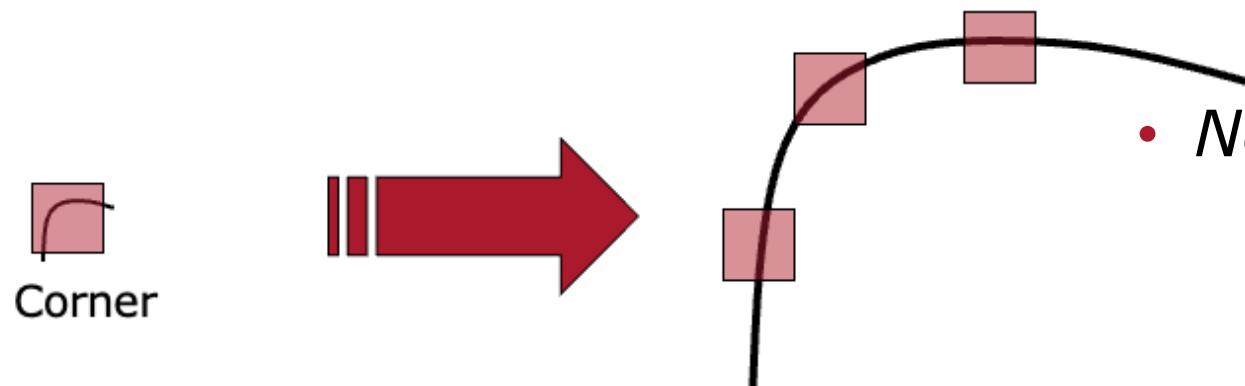
Harris corners features - Properties



- Corner location is equivariant w.r.t. translation



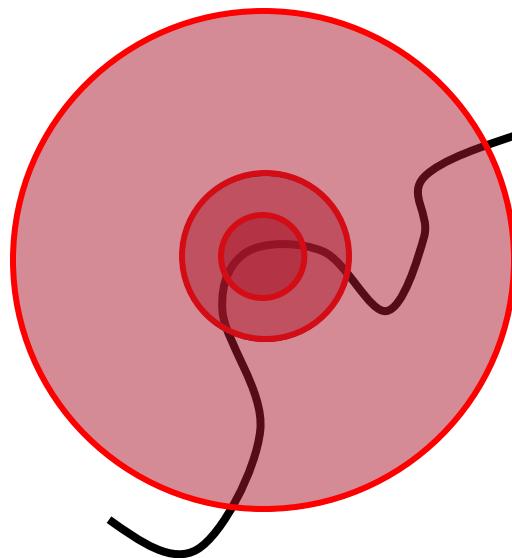
- Corner location is equivariant w.r.t. image rotation



- *No invariant* to scaling

Scale invariant detection

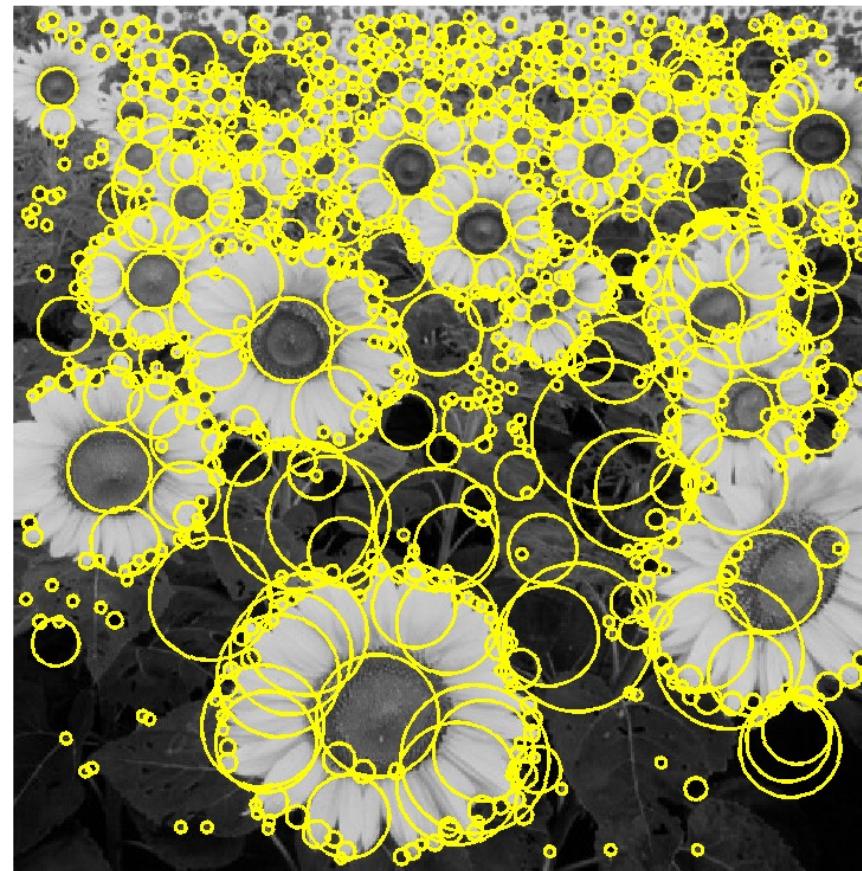
Suppose you're looking for corners



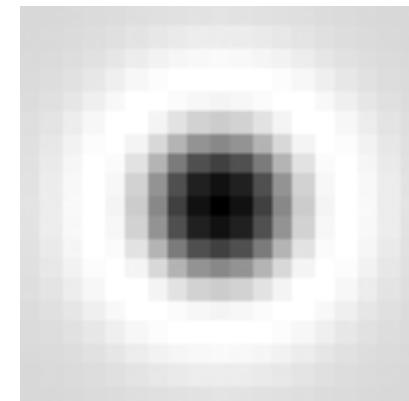
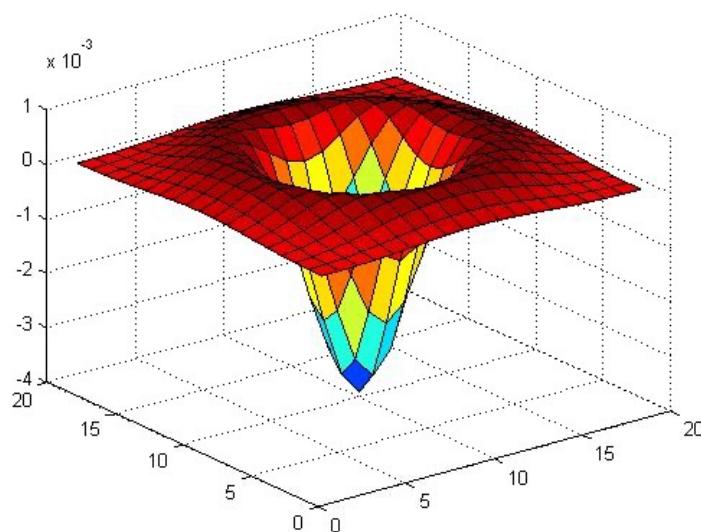
Key idea: find scale that gives local maximum of f

- in both position and scale
- One definition of f : the Harris operator

Feature extraction - Blobs



Laplacian of Gaussian (LoG)



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

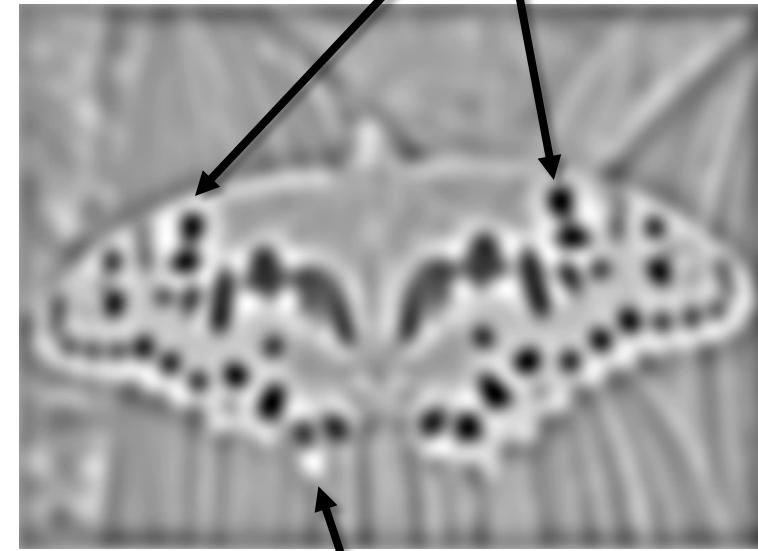
(very similar to a Difference of Gaussians (DoG) – i.e., a Gaussian minus a slightly smaller Gaussian)

Laplacian of Gaussian (LoG)

- “Blob” detector



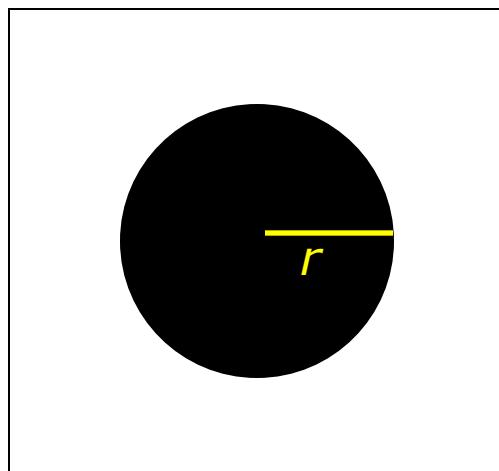
$$* \begin{array}{c} \bullet \\ \circ \end{array} =$$



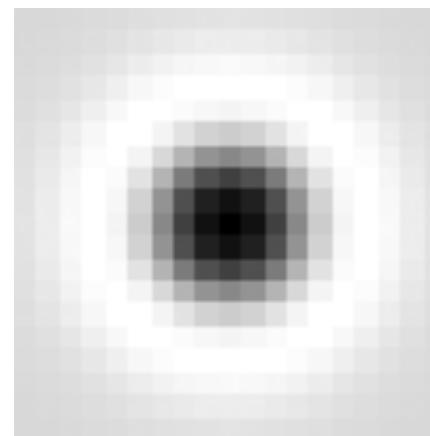
- Find maxima *and minima* of LoG operator in space and scale

Scale selection

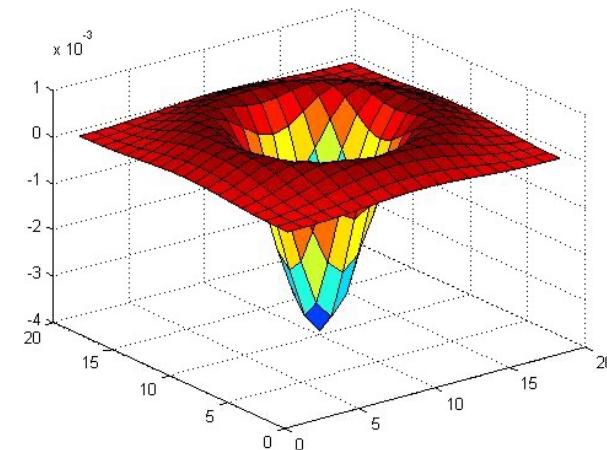
- At what scale does the Laplacian achieve a maximum response for a binary circle of radius r ?



image

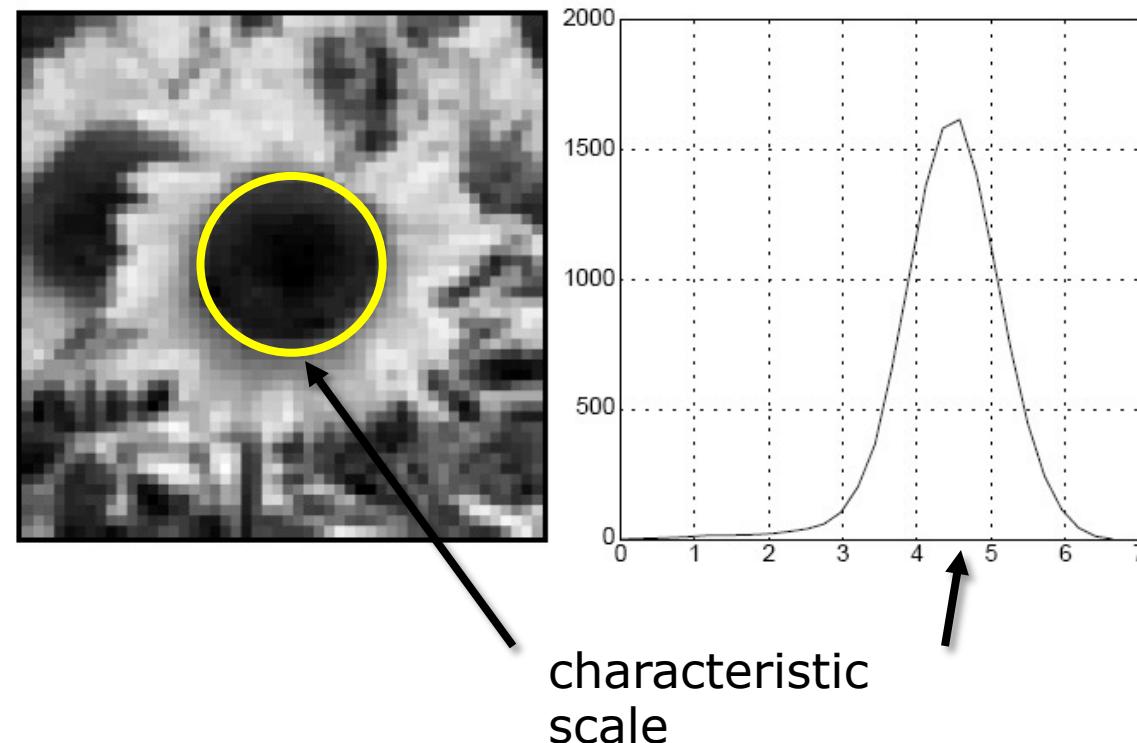


Laplacian

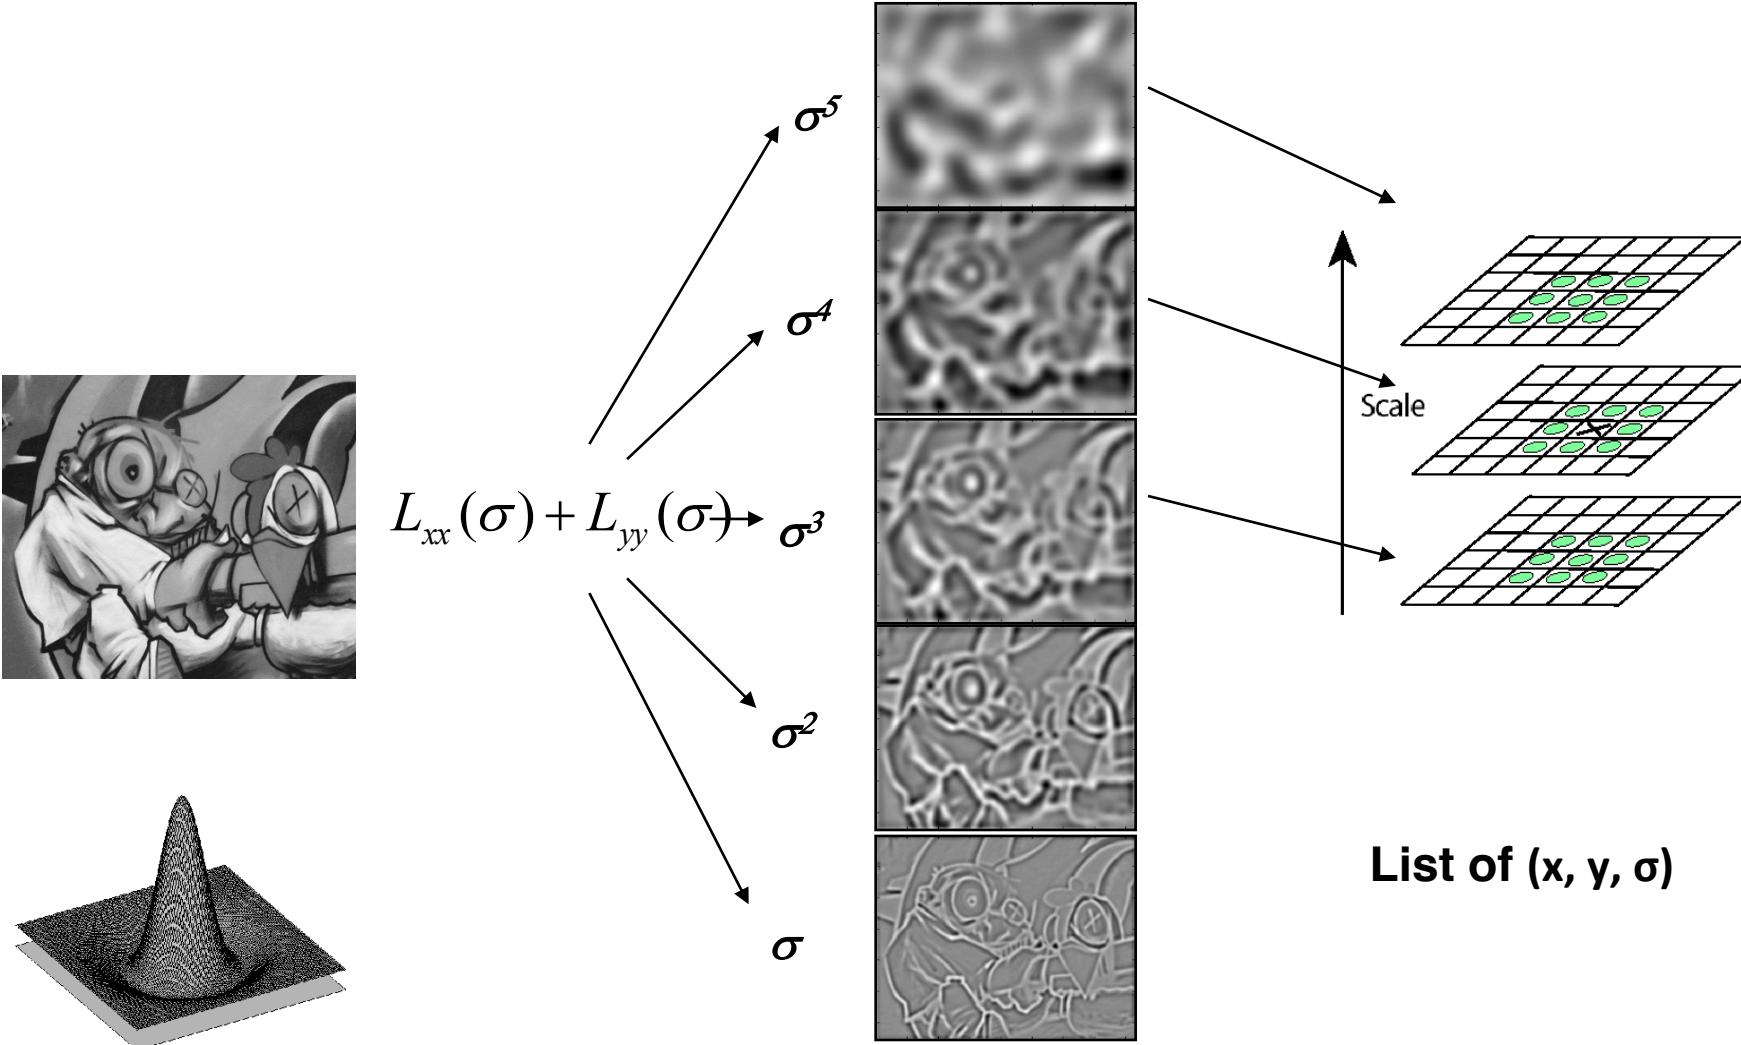


Characteristic scale

- We define the characteristic scale as the scale that produces peak of Laplacian response



Laplacian of Gaussian (LoG) Detector



Scale-space blob detector: Example



Scale-space blob detector: Example



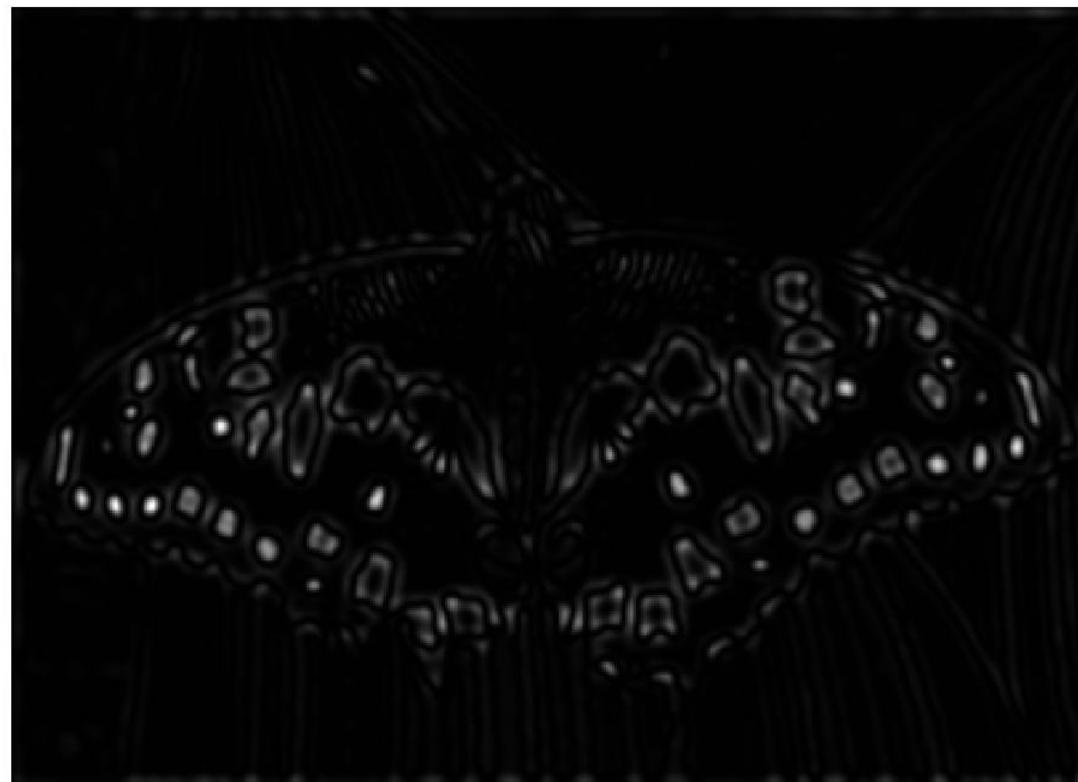
$\sigma = 2$

Scale-space blob detector: Example



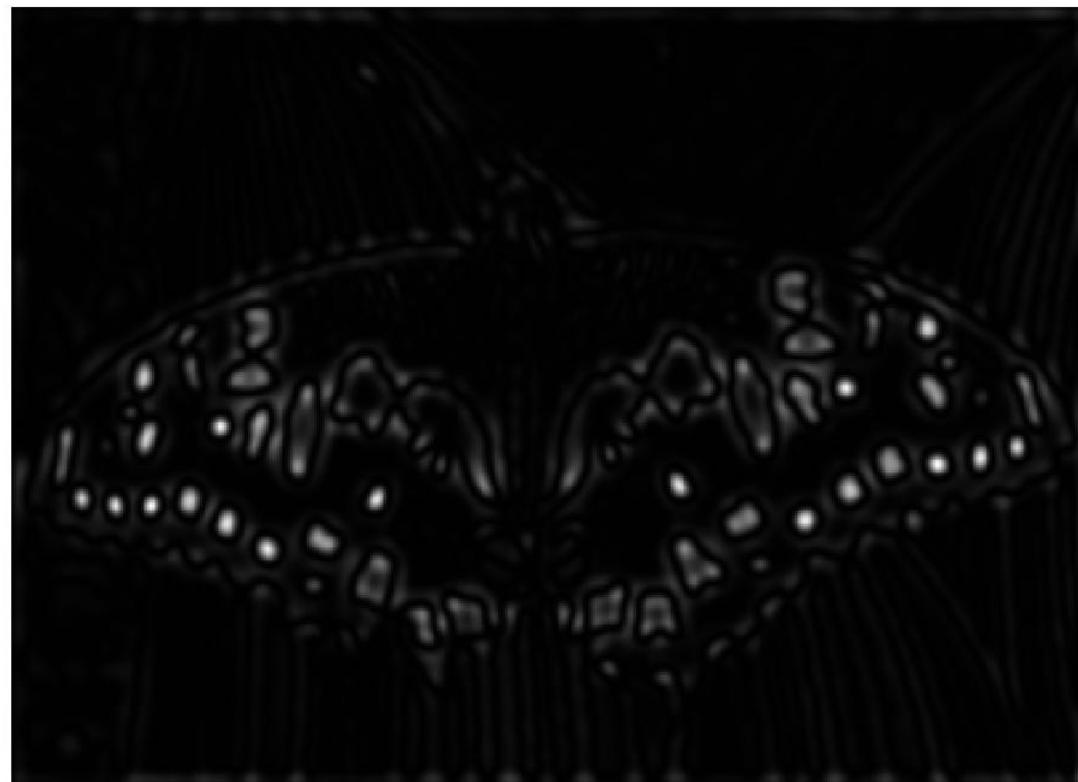
$\sigma = 2.5018$

Scale-space blob detector: Example



$\sigma = 3.1296$

Scale-space blob detector: Example



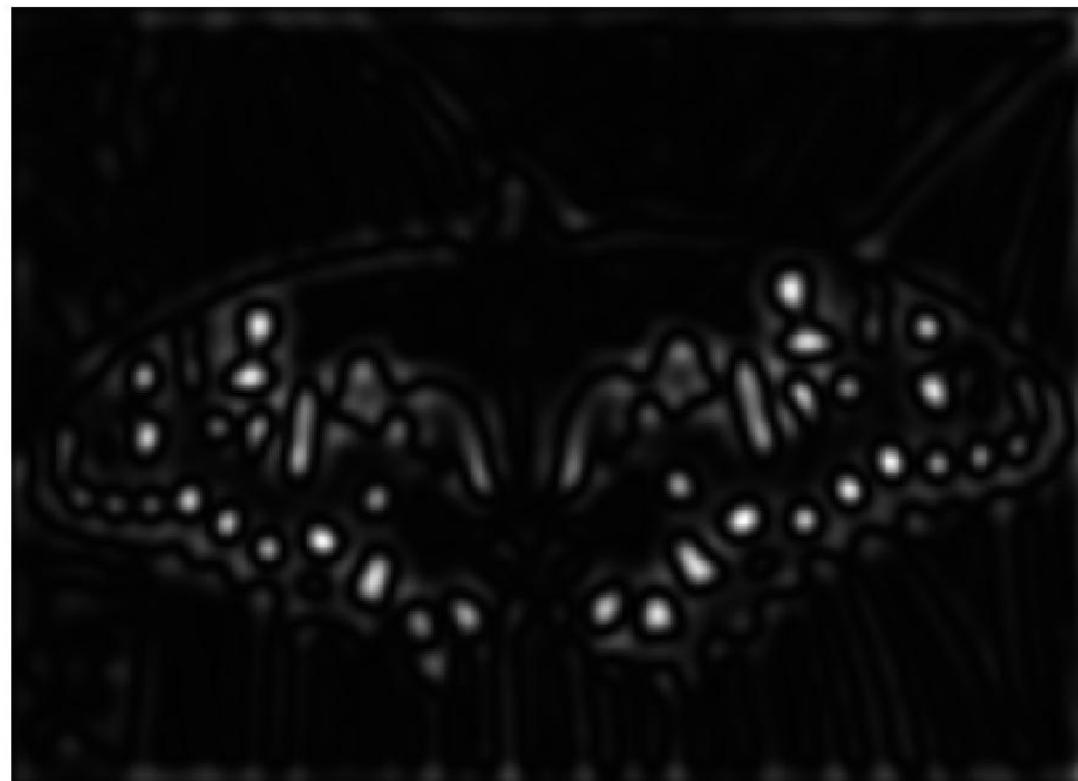
sigma = 3.9149

Scale-space blob detector: Example



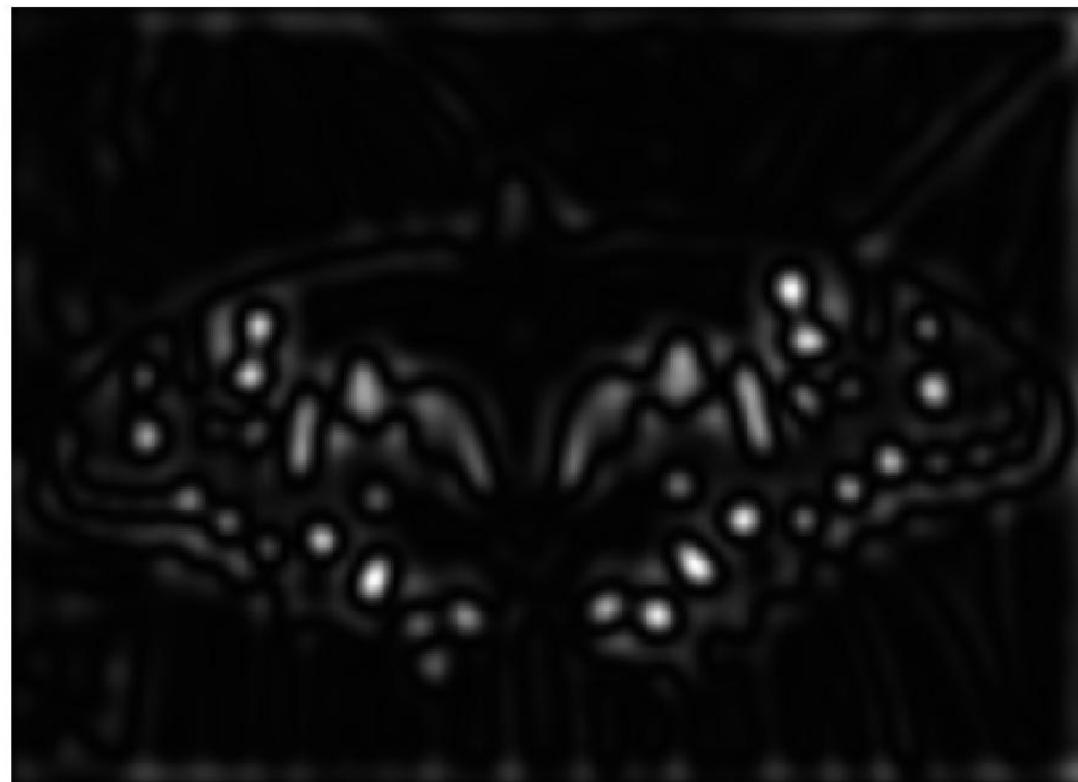
$\sigma = 4.8972$

Scale-space blob detector: Example



$\sigma = 6.126$

Scale-space blob detector: Example



$\sigma = 7.6631$

Scale-space blob detector: Example



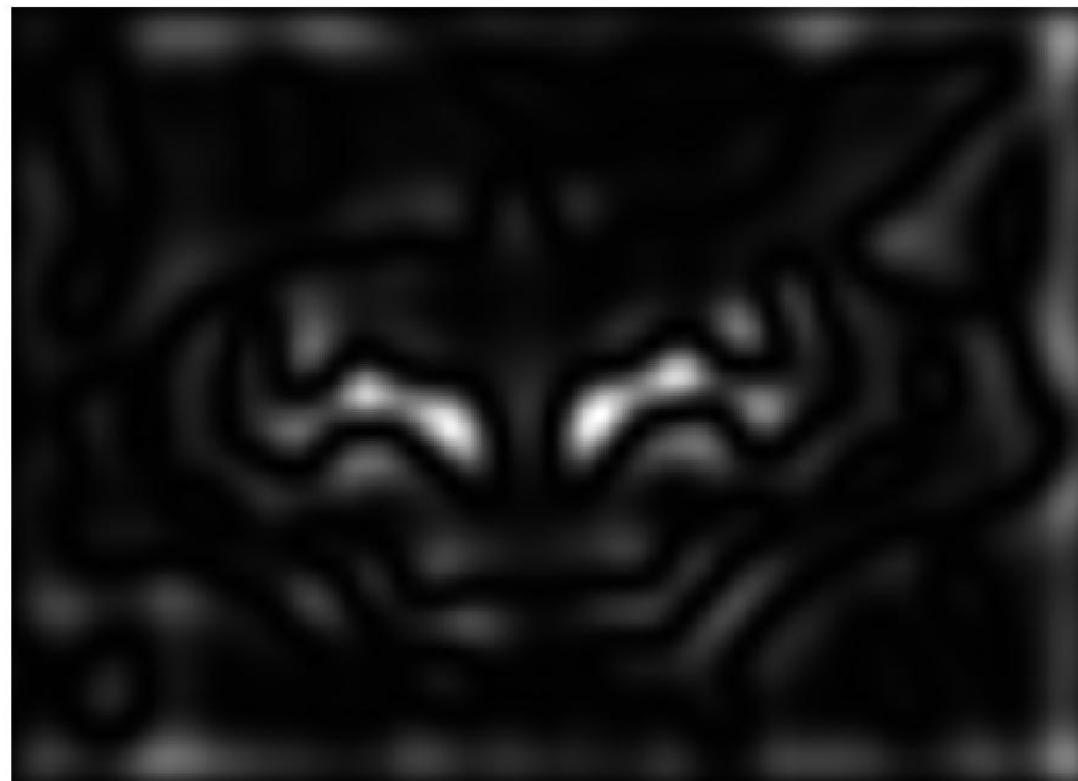
sigma = 9.5859

Scale-space blob detector: Example



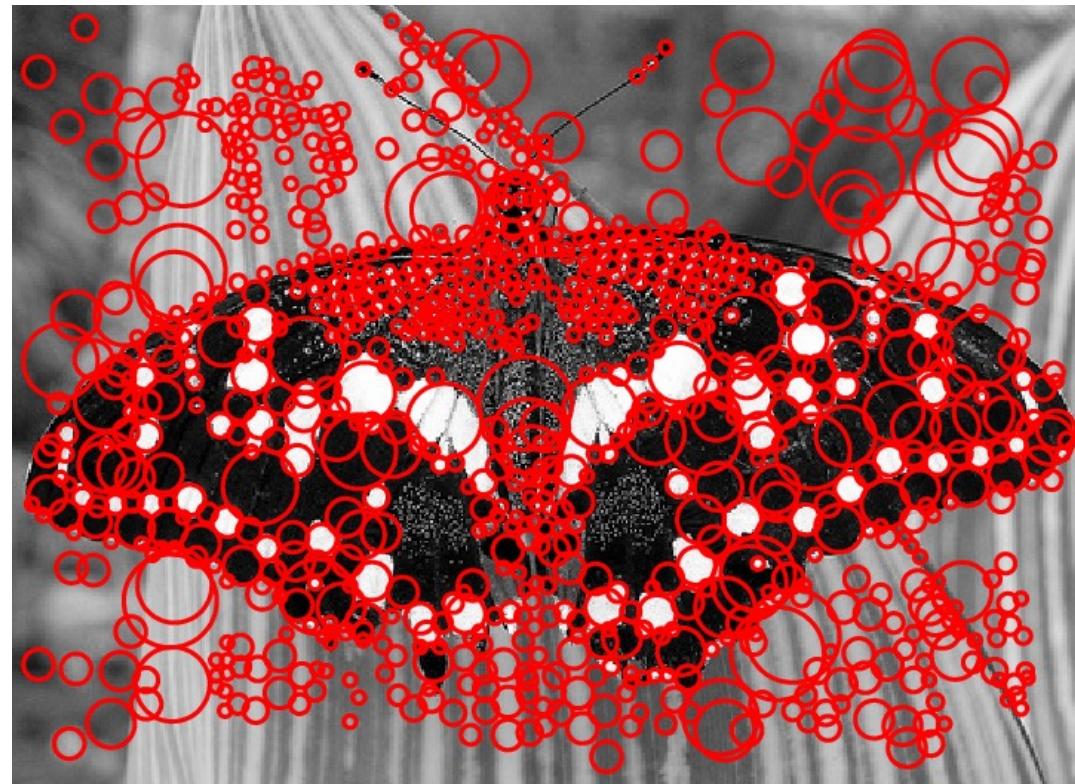
$\sigma = 11.9912$

Scale-space blob detector: Example



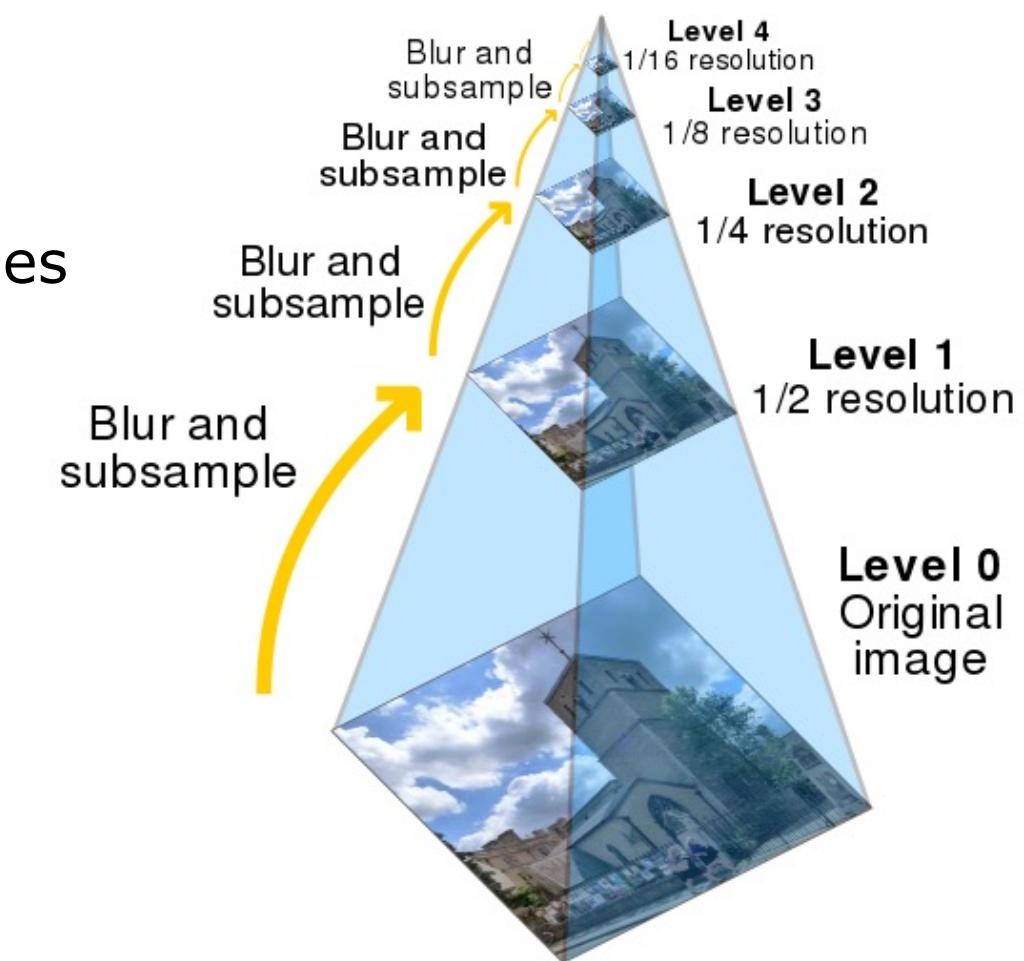
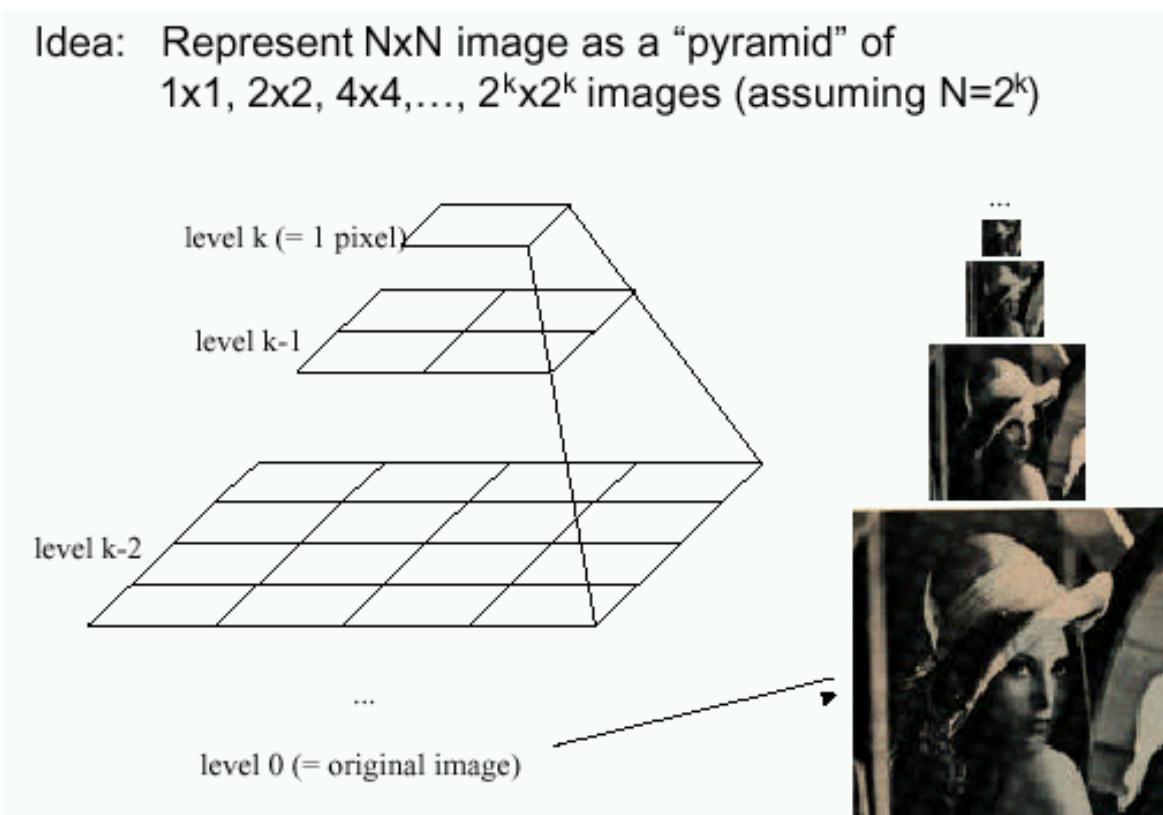
sigma = 15

Scale-space blob detector: Example



Gaussian pyramids [Burt and Adelson, 1983]

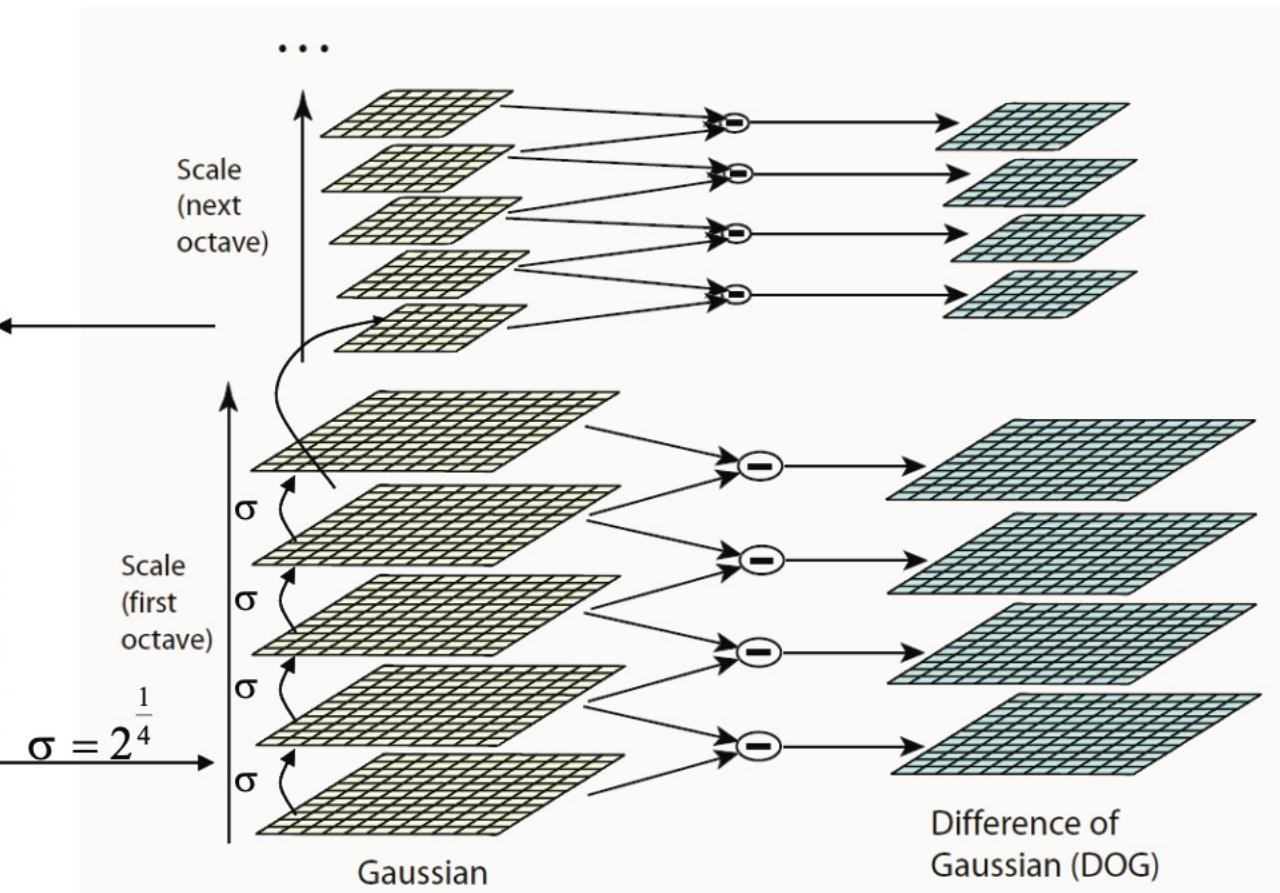
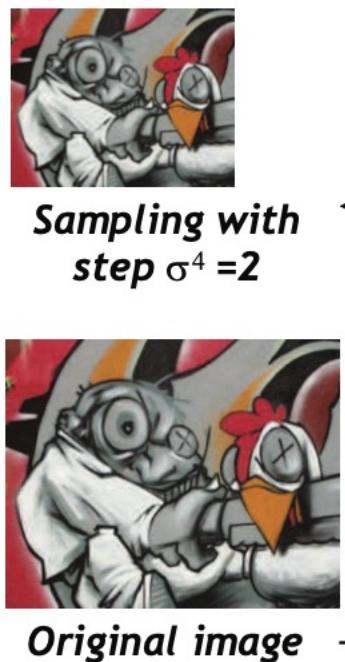
- Multi-scale representation of an image
- Gaussian smoothing and down-sampling
- Series of progressively lower-resolution images



[https://en.wikipedia.org/wiki/Pyramid_\(image_processing\)](https://en.wikipedia.org/wiki/Pyramid_(image_processing))

Difference of Gaussian (DoG) Detector

- DoG good approximator for LoG
- Subtracting Adjacent scale of a Gaussian pyramid
- Search for 3D scale space extrema



Scale Invariant Feature Transform (SIFT)

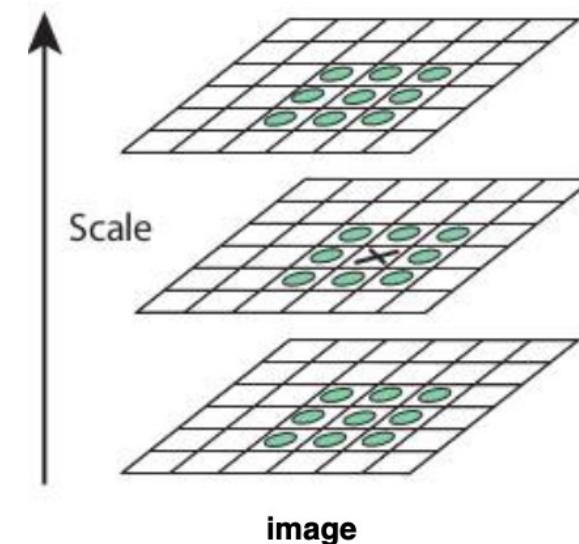
- Extract keypoints and compute its descriptors
- LoG for the detection of blobs
- Find local maxima across the scale and space → List of (x, y, σ)
- **LoG** costly → SIFT ← **DoG**
 - two different σ , let it be σ and $k\sigma$
- Different octaves of images in Gaussian pyramid
- Local extrema over scale and space
- One pixel is compared with its 8 neighbors and 9 pixels in previous scales



David Lowe

E-mail: lowe@cs.ubc.ca

**Distinctive Image Features
from Scale-Invariant Keypoints**



David G. Lowe

Computer Science Department
University of British Columbia
Vancouver, B.C., Canada
lowe@cs.ubc.ca

January 5, 2004

... end of Lecture 12 - Vision

